



Programming of Neural Networks

Today: Recurrent Neural Networks

Why do we need it?

- So far we have investigated **static** problems, that is we the shapes of the tensors in advance!
- But what if for example, we deal with text (author prediction)?

How-dy-ho! (Mr. Hankey)




Vs

Despite the constant negative press covfefe


A simple (but bad) solution

- Map every token to a vector of 100 numbers, and average it

How-dy-ho!




Despite the constant negative press covfefe



Why is this a bad solution?

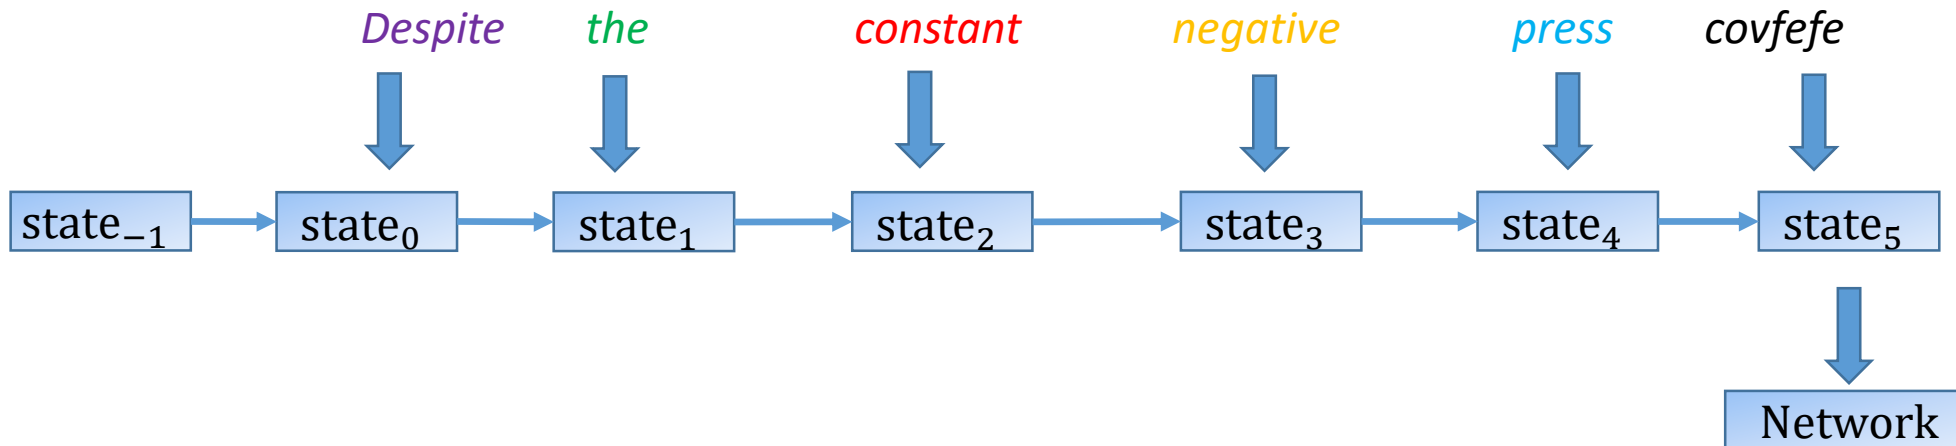
- The uniqueness of a token gets completely destroyed
- The order of the tokens is not taken into account („Bag of words“)
- The amount of tokens is ignored

➔ What would be a better solution?



Motivation: Recurrent Networks

- Instead of simply averaging the tokens, we could iteratively update a **state** and use the final state into our network




➔ We call this structure a **Recurrent Structure**

Modelling the state update

- While it is rather clear that the state can be modelled as a vector, how to update the internal state remains unclear
- ➔ How can we force a machine to compress (and store) the important information for a task in a single vector?
- Many different approaches, e.g.
 - Vanilla Recurrent Cell
 - Gated Recurrent Unit (GRU)
 - Long Short Term Memory (LSTM)

The Vanilla Recurrent Cell (one possibility)

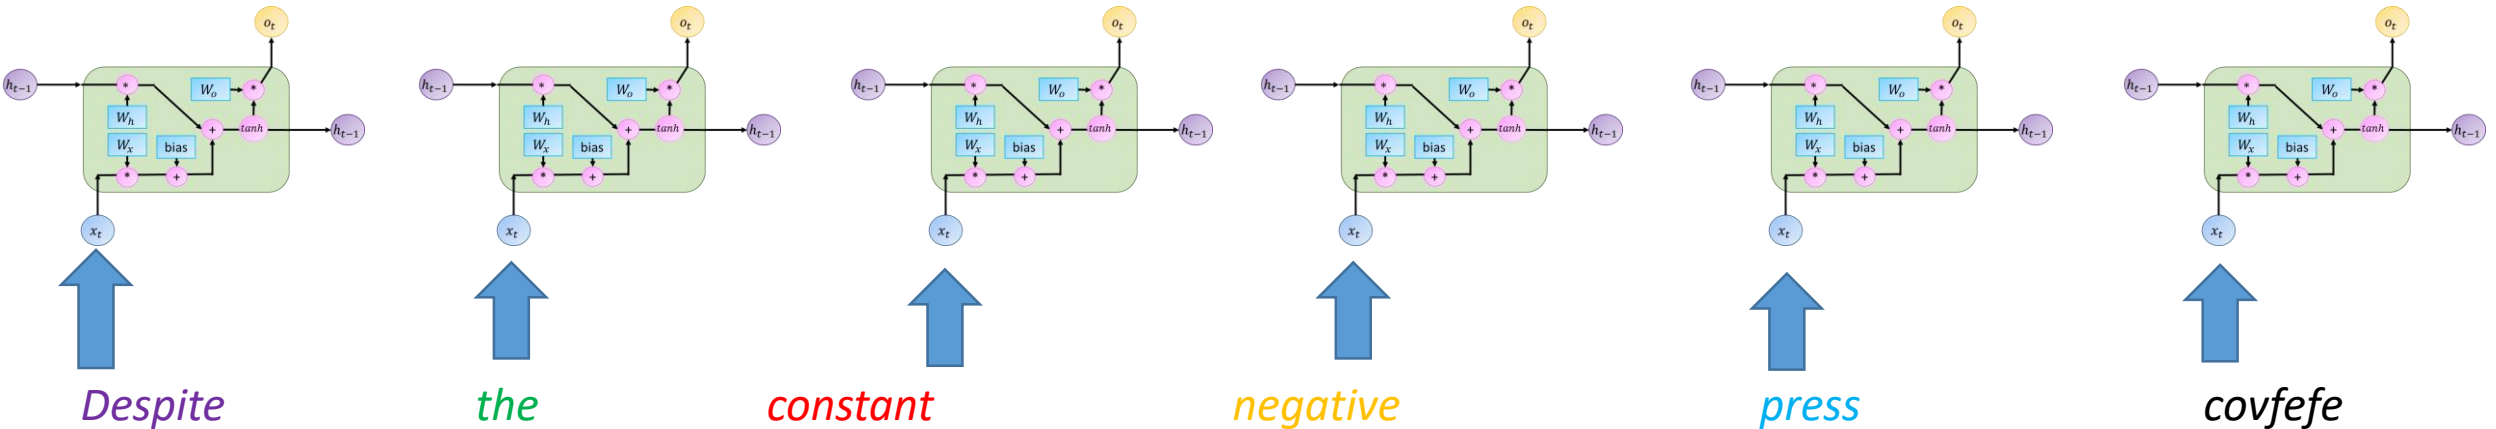


$$h_t = \tanh(h_{t-1} \cdot W_h + x_t \cdot W_x + bias)$$


$$o_t = h_t \cdot W_o$$

The Vanilla Recurrent Cell in action


We can apply the Recurrent Cell by **unrolling** (copy + pasting) it as many times as required



Vanilla Recurrent Under the hood



Vanilla Recurrent Under the hood

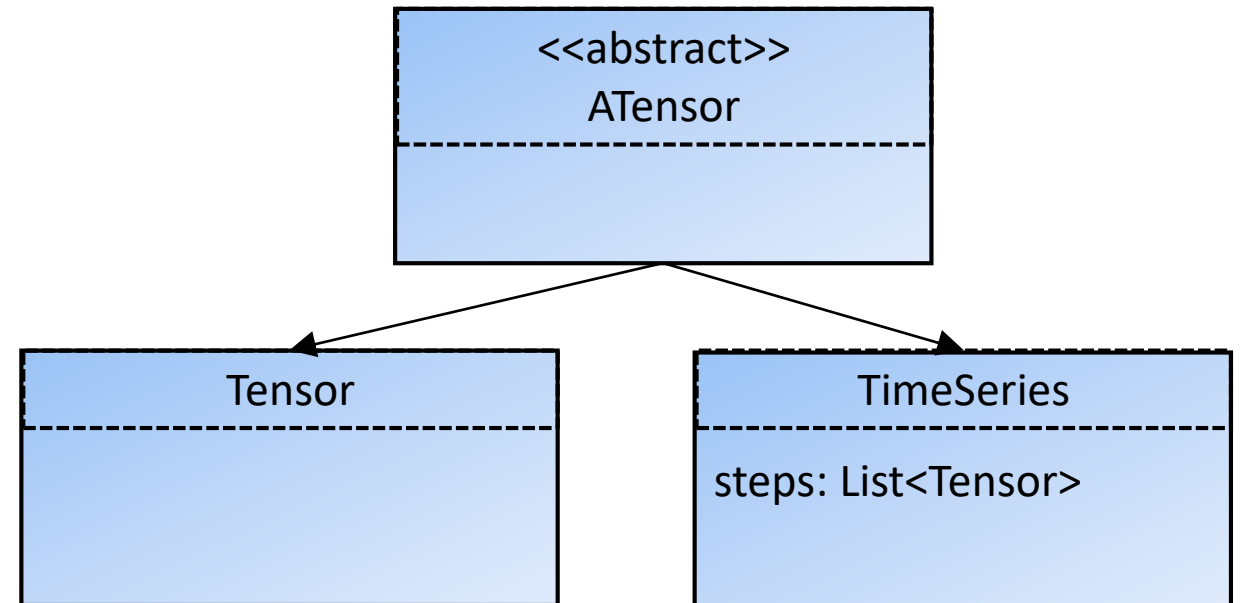


Vanilla Recurrent: The framework

- As its input it requires a „Timeseries“, which is a list of tensors that is naturally ordered by time
- Its output is either:
 - a single Tensor (just the last output)
 - A Timeseries of Tensors (a list of all outputs)

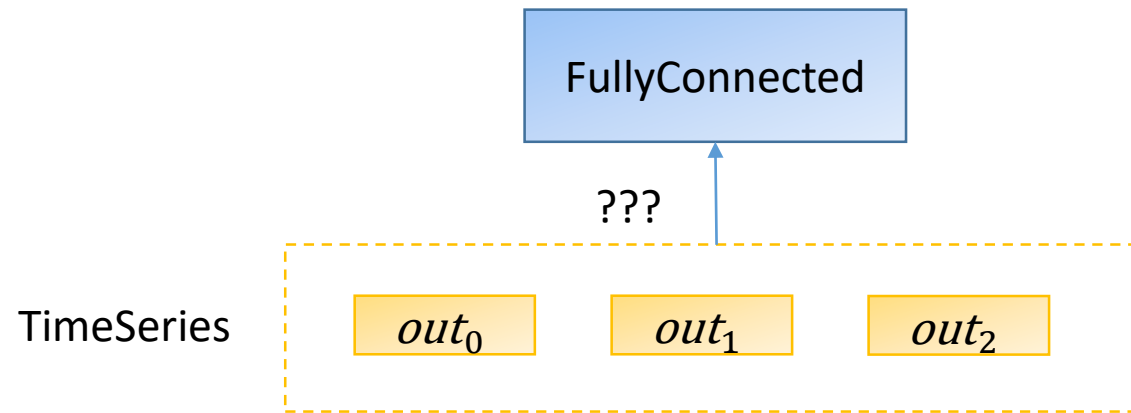
Extending the Tensor class

- A tensor is now either:
 - A regular Tensor
 - or a TimeSeries
- No further changes to the layer interfaces




Implications to the framework

- What should (for example) a Fully Connected Layer do with a TimeSeries?



Implications to the framework

➔ It is applied (with the same parameters) to **every tensor** in the time series separately, and hence produces another **TimeSeries**



Implications to the framework

- Integrating a TimeSeries into the framework, forces you to rework quite a bit of code, since every layer needs to decide what it is going to do with it!



Vanilla Recurrent: Parameters

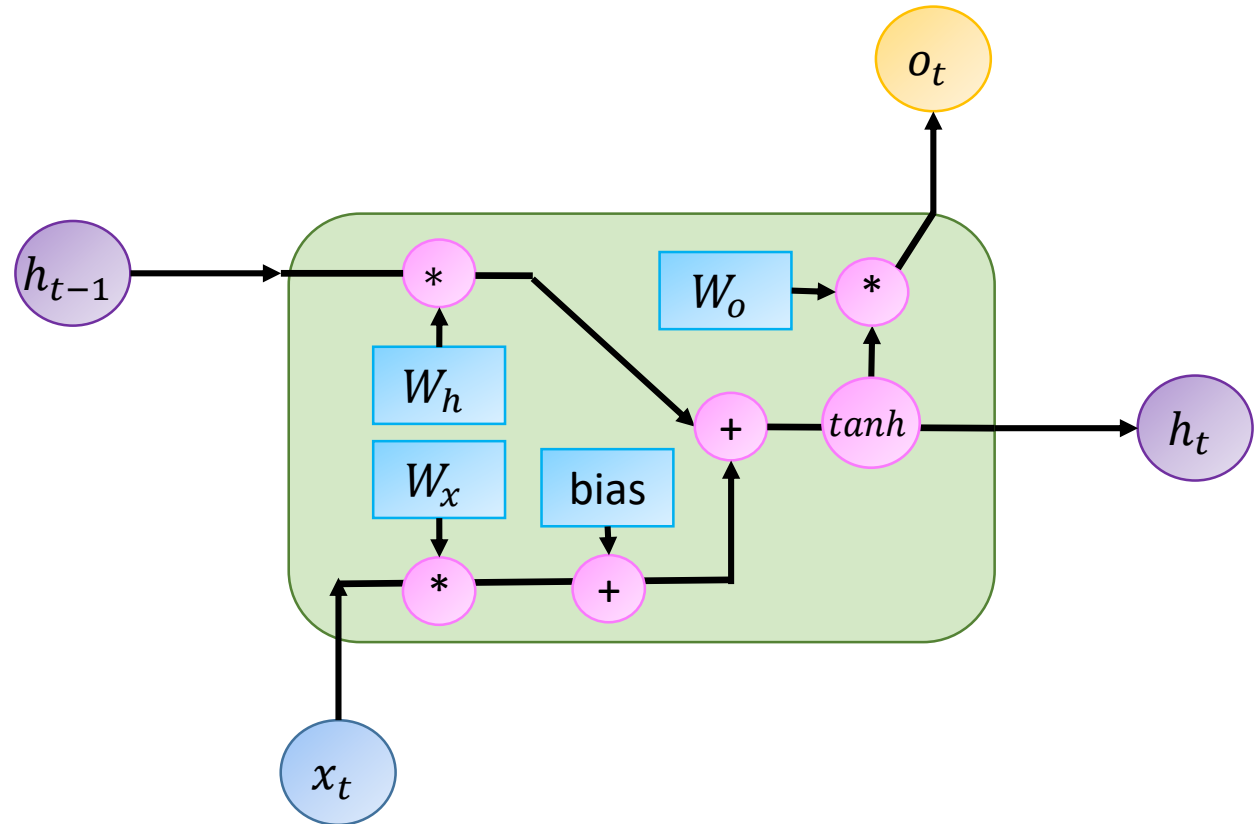
W_h : recurrent weights, shape $[h,h]$

W_x : input weights, shape $[x,h]$

W_o : output weights, shape $[h,out]$

bias:, shape $[h]$

→ Hidden states have to be stored somewhere clever, so that you got access to all of them during the backward pass!



Vanilla Recurrent: The math

- Forward pass is straight forward (pun intended)

$$\begin{aligned}g_t &= h_{t-1} \cdot W_h + x_t \cdot W_x + bias \\h_t &= \tanh(g_t) \\o_t &= h_t \cdot W_o\end{aligned}$$

- h_{-1} is initialized as a zero Tensor of according size

Vanilla Recurrent: Backward

- We have to derive 4 equations:

$$\delta o = ???$$

$$\delta h = ???$$

$$\delta g = ???$$


$$\delta x = ???$$



Why a delta for the input x ???

Vanilla Recurrent: Backward

- We have to derive 4 equations δo_t :
- δo_t : These are the deltas that are passed from above



Vanilla Recurrent: Backward


- We have to derive 4 equations δg_t :

$$h_t = \tanh(g_t)$$

$$\frac{\partial L}{\partial g_t} = \frac{\partial L}{\partial h_t} \cdot \frac{\partial h_t}{\partial g_t}$$

$$\delta g_t = \delta h_t \cdot (1 - (\tanh g_t)^2)$$

$$= \delta h_t \cdot (1 - h_t^2)$$



Vanilla Recurrent: Backward


- We have to derive 4 equations δh_t :
- h is involved in 2 equations:

$$g_t = h_{t-1} \cdot W_h + x_t \cdot W_x + bias$$

$$o_t = h_t \cdot W_o$$

$$\frac{\partial L}{\partial h_t} = \frac{\partial L}{\partial o_t} \cdot \frac{\partial o_t}{\partial h_t} + \frac{\partial L}{\partial g_{t+1}} \cdot \frac{\partial g_{t+1}}{\partial h_t}$$

$$\delta h_t = \delta o_t \cdot W_o^T + \delta g_{t+1} \cdot W_h^T$$




Vanilla Recurrent: Backward

- We have to derive 4 equations δx_t :

$$g_t = h_{t-1} \cdot W_h + x_t \cdot W_x + bias$$

$$\frac{\partial L}{\partial x} = \frac{\partial L}{\partial g_t} \cdot \frac{\partial g_t}{\partial x_t}$$

$$\delta x_t = \delta g_t \cdot W_x^T$$



Vanilla Recurrent: Backward

- Backward equations:

δo_t : from above

$$\delta h_t = \delta o_t \cdot W_o^T + \delta g_{t+1} \cdot W_h^T$$

$$\delta g_t = \delta h_t \cdot (1 - h_t^2)$$


$$\delta x_t = \delta g \cdot W_x^T$$

- h_{-1} is initialized as a zero Tensor of according size

Vanilla Recurrent: Weight updates

- We have to derive 4 equations for the weights:

$$g_t = h_{t-1} \cdot W_h + x_t \cdot W_x + \textit{bias}$$




$$\frac{\partial L}{\partial \textit{bias}} = \frac{\partial L}{\partial g_t} \cdot \frac{\partial g_t}{\partial \textit{bias}} \Rightarrow \text{weight sharing} \Rightarrow = \sum_t \frac{\partial L}{\partial g_t} \cdot \frac{\partial g_t}{\partial \textit{bias}}$$

$$\Delta \textit{bias} = \sum_t \delta g_t$$

Vanilla Recurrent: Weight updates

- We have to derive 4 equations for the weights:

$$g_t = h_{t-1} \cdot W_h + x_t \cdot W_x + bias$$




$$\frac{\partial L}{\partial W_x} = \frac{\partial L}{\partial g_t} \cdot \frac{\partial g_t}{\partial W_x} \Rightarrow \text{weight sharing} \Rightarrow \sum_t \frac{\partial L}{\partial g_t} \cdot \frac{\partial g_t}{\partial W_x}$$

$$\Delta W_x = \sum_t x_t^T \cdot \delta g_t$$

Vanilla Recurrent: Weight updates

- We have to derive 4 equations for the weights:

$$g_t = h_{t-1} \cdot W_h + x_t \cdot W_x + bias$$



$$\frac{\partial L}{\partial W_h} = \frac{\partial L}{\partial g_t} \cdot \frac{\partial g_t}{\partial W_h} \Rightarrow \text{weight sharing} \Rightarrow \sum_{t=1}^{|steps|} \frac{\partial L}{\partial g_t} \cdot \frac{\partial g_t}{\partial W_h}$$

$$\Delta W_h = \sum_{t=1} h_{t-1}^T \cdot \delta g_t$$


Vanilla Recurrent: Weight updates

- We have to derive 4 equations for the weights:

$$o_t = h_t \cdot W_o$$

$$\frac{\partial L}{\partial W_o} = \frac{\partial L}{\partial o_t} \cdot \frac{\partial p_t}{\partial W_o}$$

$$\Delta W_o = \sum_{t=0} h_t^T \cdot \delta o_t$$





Vanilla Recurrent: Weight Updates

- Weight Updates:

$$\Delta W_h = \sum_{t=1}^{|steps|} h_{t-1}^T \cdot \delta g_t$$

$$\Delta W_x = \sum_t x_t^T \cdot \delta g_t$$

$$\Delta bias = \sum_t \delta g_t$$

$$\Delta W_o = \sum_{t=0} h_t^T \cdot \delta o_t$$

Problems of the Vanilla Cell

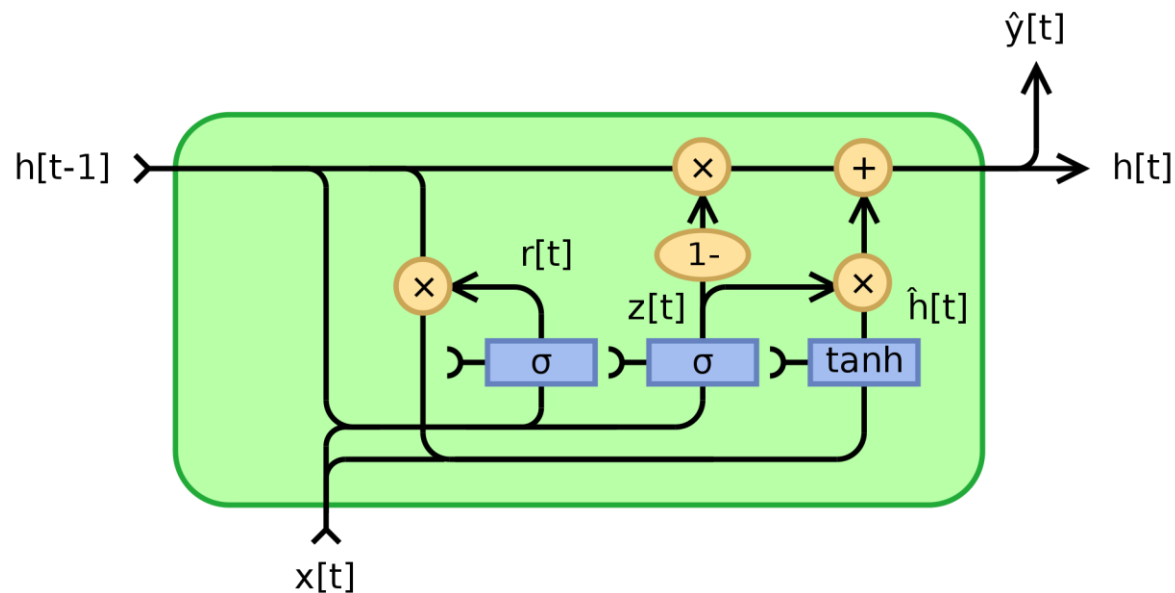
- During forward pass we did not direct the network about how it should update the internal state

$$\begin{aligned}g_t &= h_{t-1} \cdot W_h + x_t \cdot W_x + bias \\h_t &= \tanh(g_t) \\o_t &= h_t \cdot W_o\end{aligned}$$

- ➔ Can we direct it a bit so that we can integrate a controlled update and reset of the internal state?
- ➔ Also there is an issue called „Exponential Weight Decay“

Gated Recurrent Unit (GRU)

- Instead of a single gate g , it makes use of 2 gates, r and z



Source: Wikipedia: GRU

$$z'_t = h_{t-1} \cdot W_{RZ} + x_t \cdot W_Z + b_z$$

$$z_t = \sigma(z'_t)$$

$$r'_t = h_{t-1} \cdot W_{Rr} + x_t \cdot W_r + b_r$$

$$r_t = \sigma(r'_t)$$


$$f_t = r_t \cdot h_{t-1}$$

$$o'_t = f_t \cdot W_{Ro} + x_t \cdot W_o + b_o$$

$$o_t = \tanh(o'_t)$$

$$h_t = (1 - z_t) \cdot h_{t-1} + z_t \cdot o_t$$

Gated Recurrent Unit (GRU)



$$z'_t = h_{t-1} \cdot W_{RZ} + x_t \cdot W_Z + b_Z$$

$$z_t = \sigma(z'_t)$$

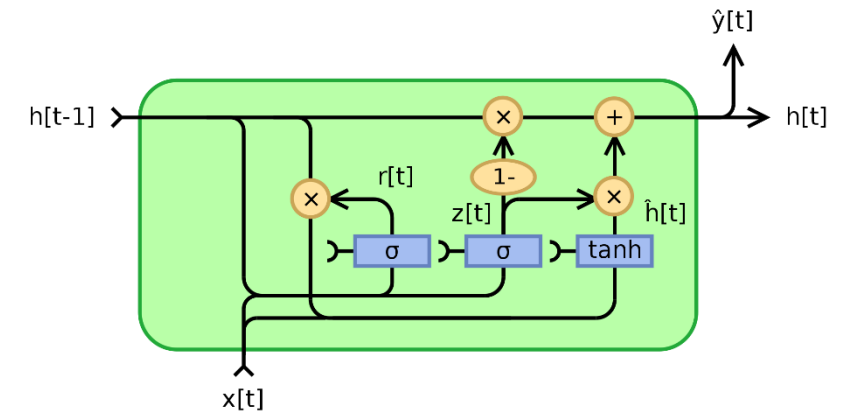
$$r'_t = h_{t-1} \cdot W_{Rr} + x_t \cdot W_r + b_r$$

$$r_t = \sigma(r'_t)$$

Input-gates: Similar to g in the vanilla state, they capture different characteristics of the input and the last state

Source: Wikipedia: GRU

Gated Recurrent Unit (GRU)



$$f_t = r_t \cdot h_{t-1}$$

$$o'_t = f_t \cdot W_{Ro} + x_t \cdot W_o + b_o$$

$$o_t = \tanh(o'_t)$$

Which entries of the last state should remain and which should be forgotten

Ratio of „What to keep“ and „What to add“

$$h_t = (1 - z_t) \cdot h_{t-1} + z_t \cdot o_t$$

Source: Wikipedia: GRU




Gated Recurrent Unit (GRU): Backward

- Your exercise! Should be fairly easy by now

Long Short Term Memory (LSTM)

- Instead of 2 Gates, involved in the GRU, it uses 3:

- **Forget f**
- **Output o**
- **Input i**



Source: Colah's Blog

$$f'_t = h_{t-1} \cdot W_{Rf} + x_t \cdot W_f + b_f$$

$$f_t = \sigma(f'_t)$$

$$o'_t = h_{t-1} \cdot W_{Ro} + x_t \cdot W_o + b_o$$

$$o_t = \sigma(o'_t)$$

$$i'_t = h_{t-1} \cdot W_{Ri} + x_t \cdot W_i + b_i$$


$$i_t = \sigma(i'_t)$$

$$a'_t = h_{t-1} \cdot W_{Ra} + x_t \cdot W_a + b_a$$

$$a = \tanh(a'_t)$$

Long Short Term Memory (LSTM)

- It also introduces an additional internal state c_t which is also updated in every time step



Source: Colah's Blog

$$c_t = a_t \cdot i_t + f_t \cdot c_{t-1}$$

$$h_t = \tanh(c_t) \cdot o_t$$

$$out_t = h_t$$



LSTM: Forward

$$a'_t = h_{t-1} \cdot W_{Ra} + x_t \cdot W_a + b_a$$

$$a = \tanh(a'_t)$$

$$f'_t = h_{t-1} \cdot W_{Rf} + x_t \cdot W_f + b_f$$

$$f_t = \sigma(f'_t)$$

$$o'_t = h_{t-1} \cdot W_{Ro} + x_t \cdot W_o + b_o$$

$$o_t = \sigma(o'_t)$$

$$i'_t = h_{t-1} \cdot W_{Ri} + x_t \cdot W_i + b_i$$

$$i_t = \sigma(i'_t)$$

$$c_t = a_t \cdot i_t + f_t \cdot c_{t-1}$$

$$h_t = \tanh(c_t) \cdot o_t$$

$$out_t = h_t$$

LSTM: Backward (Derivation on the board)

$$\delta h_t = \delta a'_{t+1} \cdot W_{Ra}^T + \delta i'_{t+1} \cdot W_{Ri}^T + \delta o'_{t+1} \cdot W_{Ro}^T + \delta f'_{t+1} \cdot W_{Rf}^T + \delta out_t$$

$$\delta c_t = \delta h_t \cdot o_t \cdot (1 - \tanh(c_t)) + \delta c_{t+1} \cdot f_{t+1}$$

$$\delta a_t = \delta c_t \cdot i_t$$

$$\delta i_t = \delta c_t \cdot a_t$$

$$\delta f_t = \delta c_t \cdot c_{t-1}$$

$$\delta o_t = \delta h_t \cdot \tanh(c_t)$$

$$\delta a'_t = \delta a_t \cdot (1 - a_t^2)$$

$$\delta i'_t = \delta i_t \cdot i_t(1 - i_t)$$

$$\delta f'_t = \delta f_t \cdot f_t(1 - f_t)$$

$$\delta o'_t = \delta o_t \cdot o_t(1 - o_t)$$

$$\delta x_t = \delta a'_t \cdot W_a^T + \delta i'_t \cdot W_i^T + \delta o'_t \cdot W_o^T + \delta f'_t \cdot W_f^T$$



LSTM: Weight Updates

Bias

$$\Delta b_a = \sum_t \delta a'_t$$

$$\Delta b_i = \sum_t \delta i'_t$$

$$\Delta b_f = \sum_t \delta f'_t$$

$$\Delta b_o = \sum_t \delta o'_t$$

Input-Weights

$$\Delta W_a = \sum_t x_t^T \cdot \delta a'_t$$

$$\Delta W_i = \sum_t x_t^T \cdot \delta i'_t$$

$$\Delta W_f = \sum_t x_t^T \cdot \delta f'_t$$

$$\Delta W_o = \sum_t x_t^T \cdot \delta o'_t$$

Recurrent-Weights

$$\Delta W_{Ra} = \sum_{t=1} h_{t-1}^T \cdot \delta a'_t$$

$$\Delta W_{Ri} = \sum_{t=1} h_{t-1}^T \cdot \delta i'_t$$

$$\Delta W_{Rf} = \sum_{t=1} h_{t-1}^T \cdot \delta f'_t$$

$$\Delta W_{Ro} = \sum_{t=1} h_{t-1}^T \cdot \delta o'_t$$

Beyond LSTMs

- In recent years, there was critique towards LSTMs , that they (in theory) are capable of capturing information for a long period, but in reality they dont
- The sequential structure can further be exploited by:
 - ➔ Introduction of the „Attention Mechanism“
 - ➔ Introduction of **Highway Networks**, especially **Recurrent Highway Networks (RHN)**
 - ➔ **Conditional Random Fields** to make use of the structure of the classification label as well
 - ➔ **Mogrifier LSTM**

Highway Networks

- Highway Networks were introduced in order to ease the training of very deep networks (>100 layer)



- They are very simple in nature and can be considered a replacement for Fully Connected Layers, when training Deep Networks


Highway Networks

- Instead of just forwarding $y = x \cdot W + b$ as in Fully Connected layers
- The forward equation is changed to:

$$y = (x \cdot W + b) \cdot T(x) + x \cdot (1 - T(x))$$

With:

$$T(x) = \sigma(x \cdot W_T + b_T)$$



Highway Networks

- Let us introduce \hat{y} to be the forward variable we get from a standard Fully Connected Layer

$$y = \hat{y} \cdot T(x) + x \cdot (1 - T(x))$$

- Since T uses a sigmoid its output is in **[0,1]**

- If $T(x)[i]$ is 0 $\rightarrow y = x$

- If $T(x)[i]$ is 1 $\rightarrow y = \hat{y}$

- Else it is a mixture between both!

The signal can simply pass forward in a very quick (High)way

Highway Networks

- You have to pay attention to the shapes! Otherwise the addition does not work !

$$y = \hat{y} \cdot T(x) + x \cdot (1 - T(x))$$



Recurrent Highway Networks (RHN)

- Are the natural extension of a Highway Layer into a Recurrent Form
 - The Highway Layers come with their built-in gating mechanisms anyway!
 - Use **Microcells** and stack these before forwarding to next timestep
- (Note: Bias and activation ignored in the diagram!)

Recurrent Highway Networks (RHN)-Microcell

$$s_t^l = h_t^l \cdot t_t^l + s_{t-1}^{l-1} \cdot c_t^l$$


With:

$$h_t^l = \tanh(x_t \cdot W_H \cdot \mathbf{1}_{l=0} + s_{l-1} \cdot W_{RHl} + b_{Hl})$$

$$t_t^l = \sigma(x_t \cdot W_T \cdot \mathbf{1}_{l=0} + s_{l-1} \cdot W_{Rtl} + b_{Tl})$$


$$c_t^l = \sigma(x_t \cdot W_C \cdot \mathbf{1}_{l=0} + s_{l-1} \cdot W_{Rcl} + b_{Cl})$$

Authors recommend: $c_t^l = (1 - t_t^l)$



Recurrent Highway Layer – Recurrent Cell

Every depth l has
different weight
matrices!



Recurrent Networks -Summary

- Enables a Neural Network to capture and compress important information of a sequence into a state.
- No clear evidence which Recurrent Cells are the best, but LSTMs are the de facto standard.
- Also there are many variations of the presented cells (e.g. peephole LSTMs, etc...)
- It becomes evident, that manually deriving the backward equations is cumbersome! ➔ Next lecture: **Reverse Mode Automatic Differentiation**