

# Image Segmentation

Overview and Fully Convolutional Networks



# Object localization and detection

**Classification**



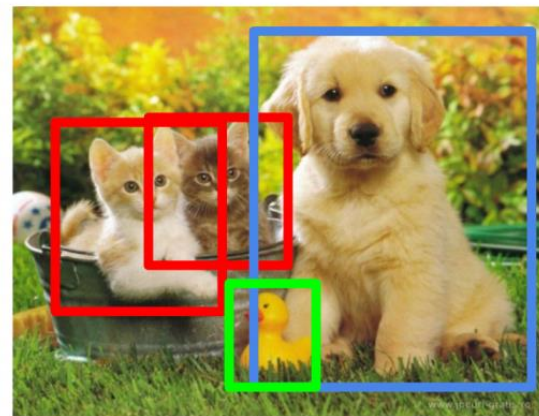
CAT

**Classification  
+ Localization**



CAT

**Object Detection**



CAT, DOG, DUCK

**Instance  
Segmentation**



CAT, DOG, DUCK

Single object

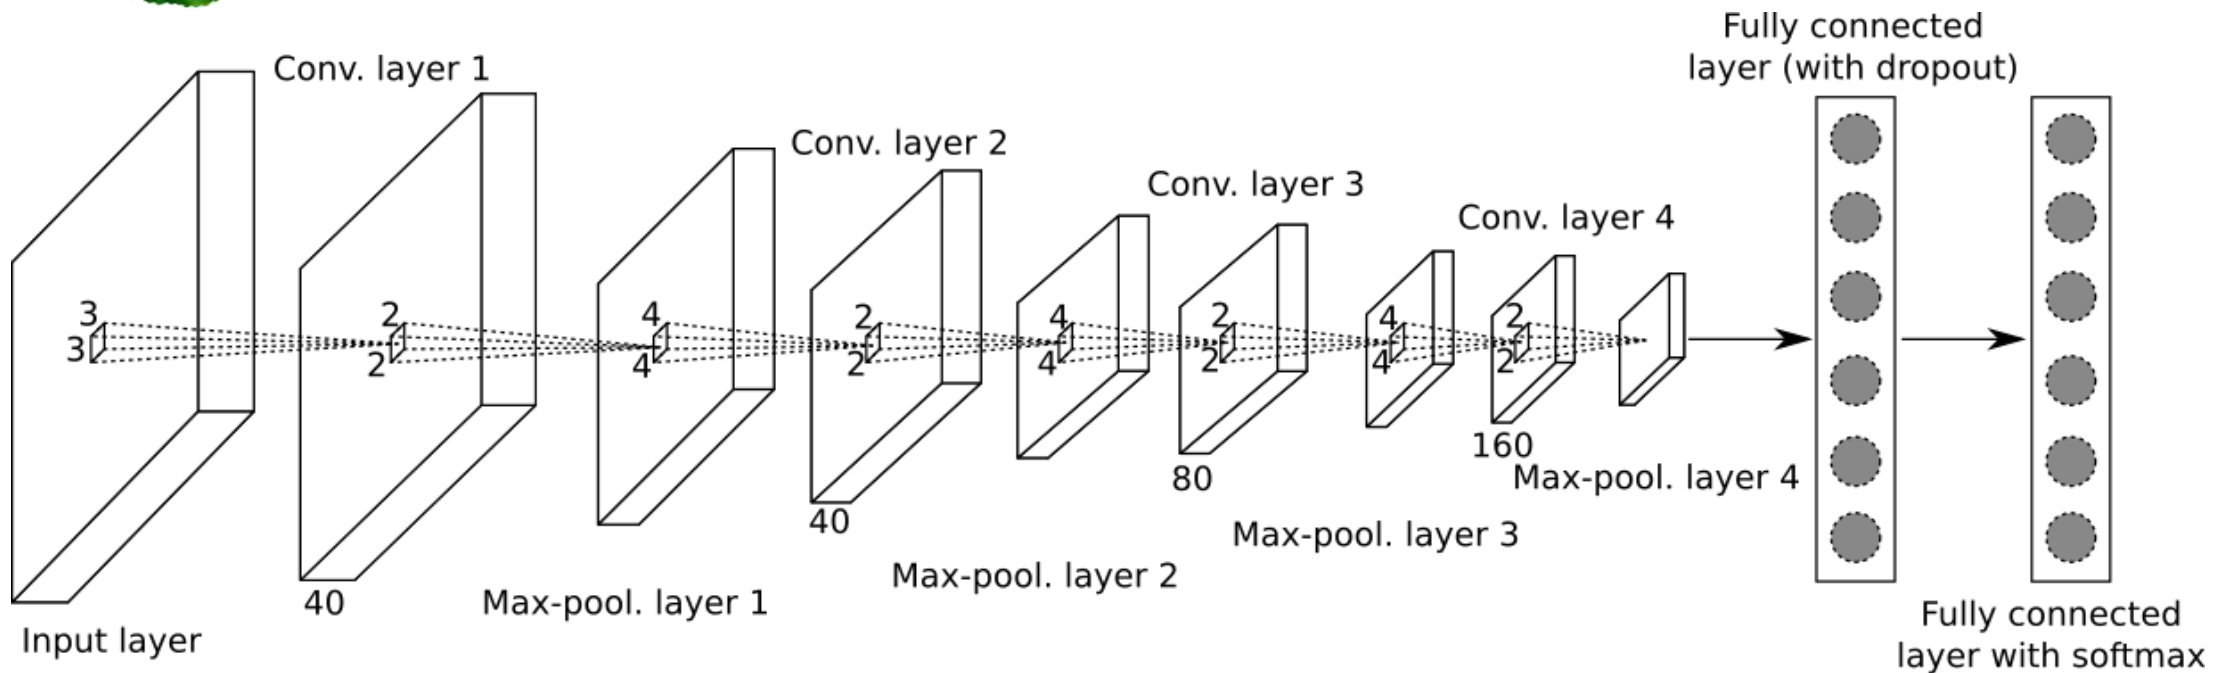
Multiple objects



# Recap: Image Classification



→ Class



# Image segmentation

Goal:

- Identify a class for every image point/pixel
- Input: Image of arbitrary dimensions and size
- Output: segmentation mask



# Image segmentation: Example

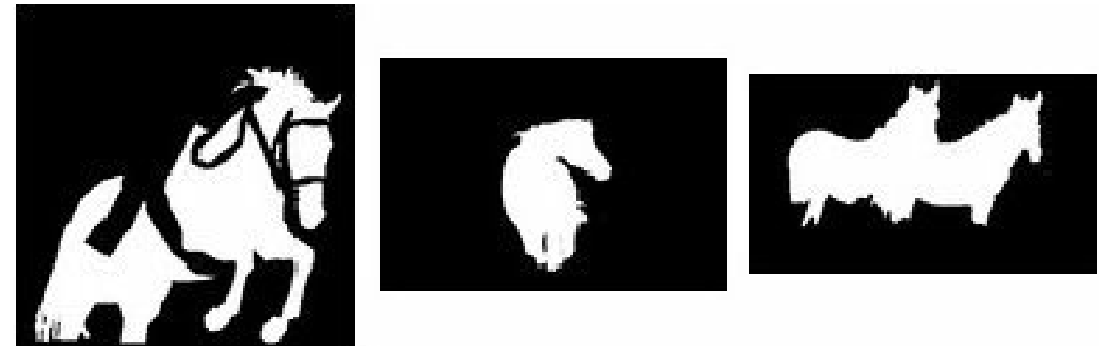
## Binary classification

Input



$$h \times w \times c$$

Output



$$h \times w \times 1$$



# Image segmentation: Example

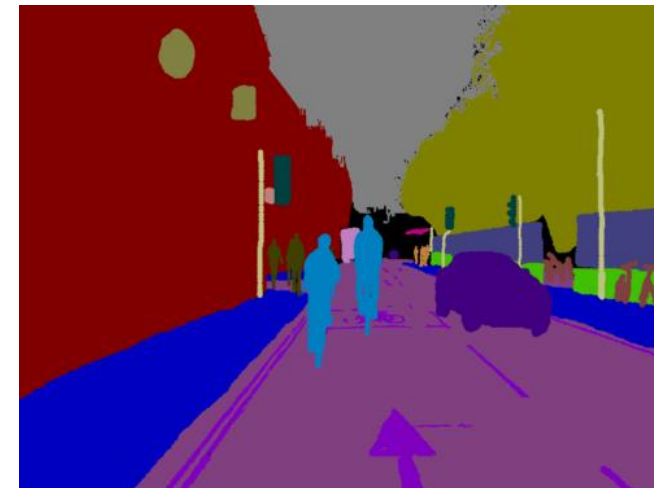
Classification with several classes (non-mutual as well)

Input



$$h \times w \times c$$

Output (mapped)



$$h \times w \times n_{\text{classes}}$$



# Table of Contents

- Sliding Window
- Fully Convolutional Network
  - Up-Sampling variations
  - Training and evaluation
  - Datasets/Transfer-Learning
- Instance Segmentation





# Sliding-Window



## Principle

- Generate many small image patches (e.g. as a grid)
- Use a CNN to classify each one of them separately
- Center pixel is the target class





# Sliding-Window (200x200 with Stride 100)



Background

Traffic  
light

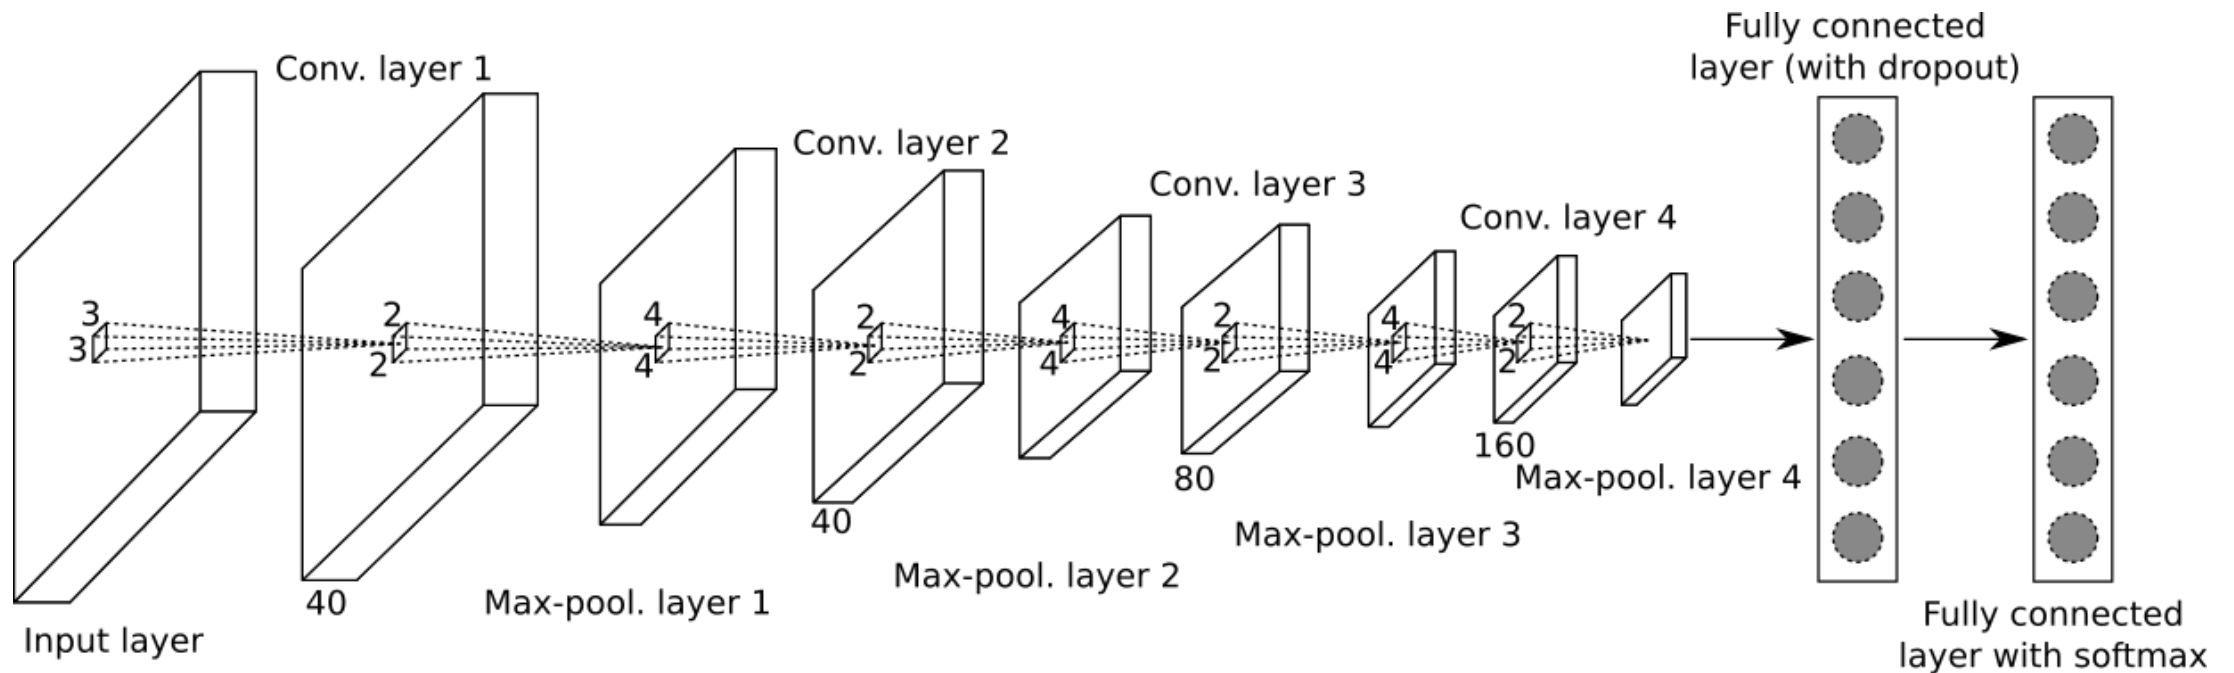
...



# Sliding-Window



→ Traffic light



# Sliding-Window: Example Page segmentation

Conv, 40, 3x3, padding=same

ReLU

MaxPool, 2x2

Conv, 40, 3x3, padding=same

ReLU

MaxPool, 2x2

Conv, 80, 3x3, padding=same

ReLU

MaxPool, 2x2

Dense(100)

ReLU

Dropout(0.5)

Dense(n\_classes)

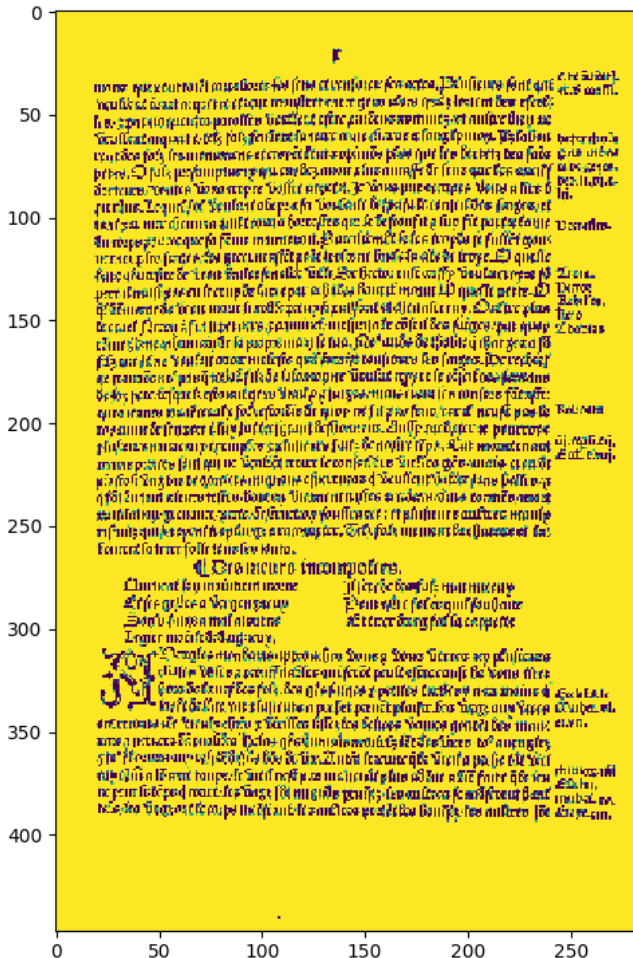
Softmax

```
model = Sequential([
    Conv2D(40, (3, 3), padding='same', activation='relu')
    MaxPooling2D(pool_size=(2, 2), strides=2),
    Conv2D(40, (4, 4), padding='same', activation='relu'),
    MaxPooling2D(pool_size=(2, 2), strides=2),
    Conv2D(80, (4, 4), padding='same', activation='relu'),
    MaxPooling2D(pool_size=(2, 2), strides=2),
    Flatten(),
    Dense(100, activation='relu'),
    Dropout(dropout),
    Dense(num_classes, activation='softmax', name='softmax'),
])
```

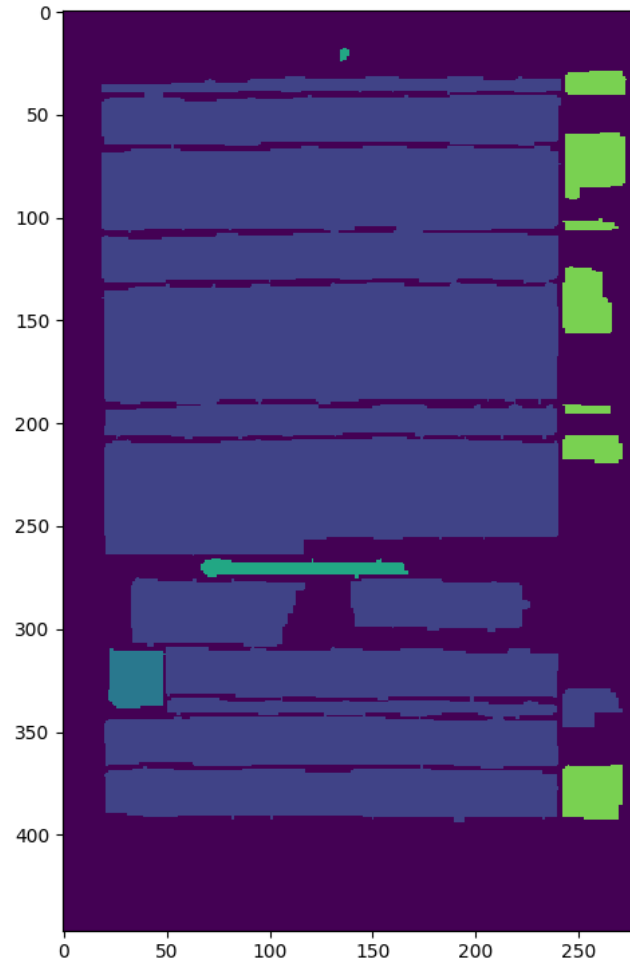


# Sliding-Window: Example Page segmentation

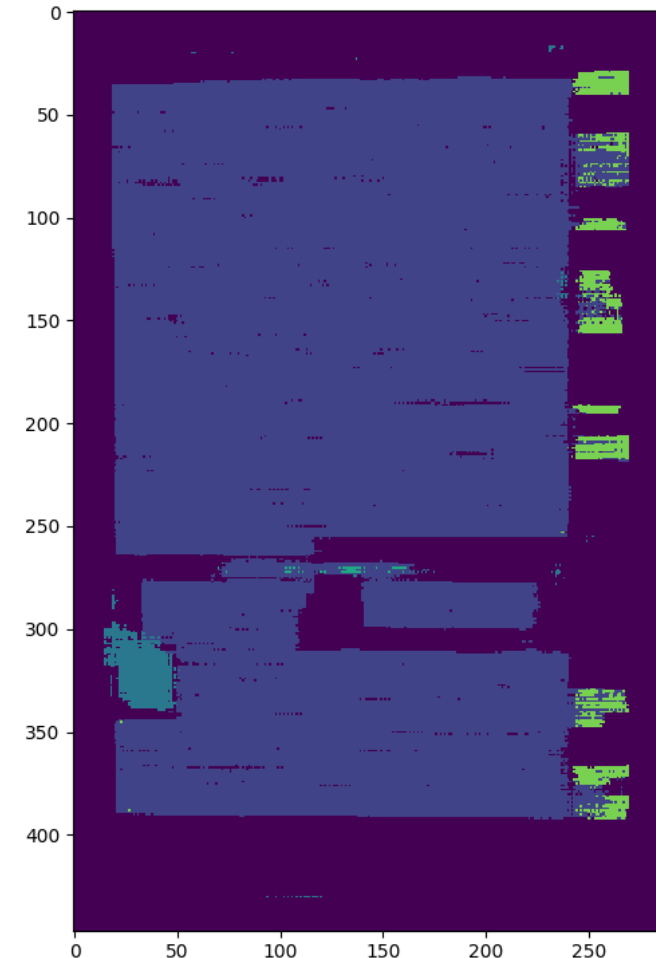
Input (Binary image)



Ground Truth (Label image)



Prediction



# Sliding-Window

## Problems

- Small surfaces:
  - Neighboring predictions have no influence on each other
- Resolution:
  - Choosing Sliding window size and stride → Scaling of original image
  - Input and output do not have matching dimensions (only with stride 1 and padding)
- Multiple computations:
  - With small stride many image patches will be sent through the CNN several times
  - Repetition of the same calculations (with the exception of the FC-Layer)



# Fully Convolutional Networks



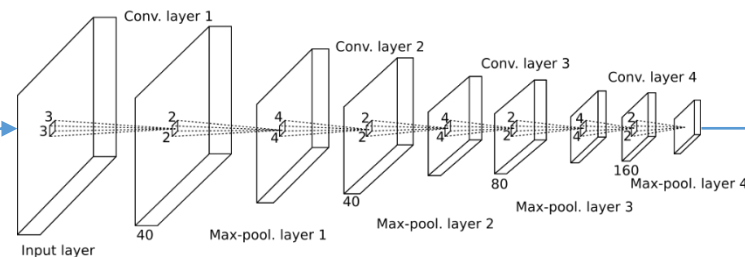


# Fully Convolutional Network for segmentation



Entire image as input

$$h \times w \times 3$$



CNN computes  
„compressed form“  
**Encoder**

$$\frac{h}{16} \times \frac{w}{16} \times k$$



Reversing pooling

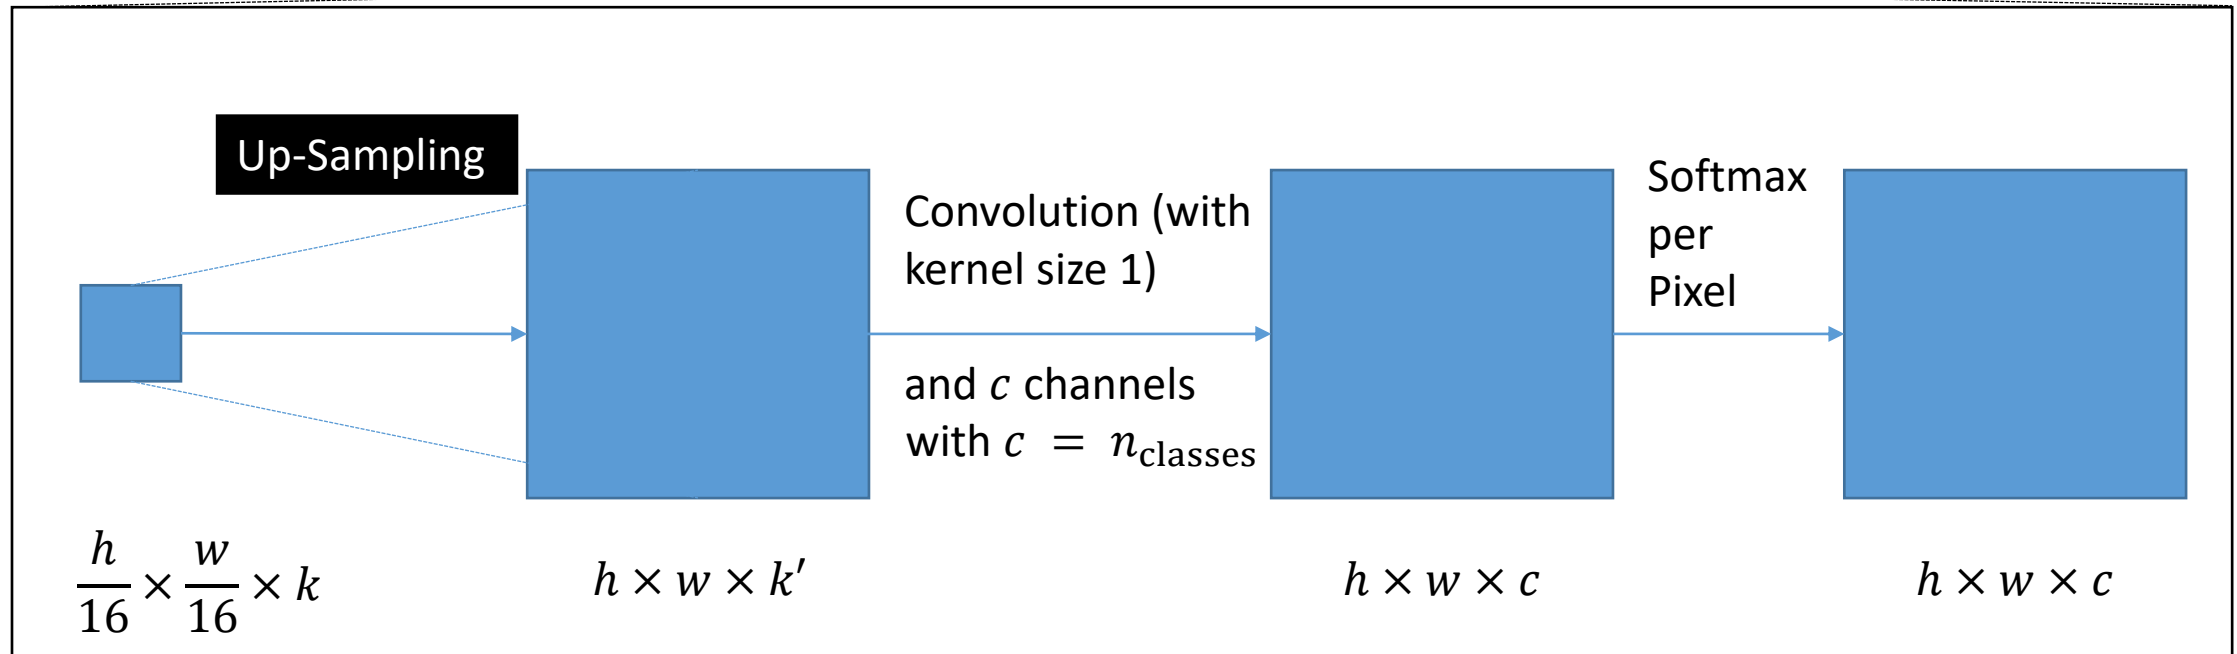
**Decoder**



$$h \times w \times n_{\text{classes}}$$

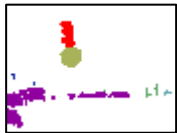


# Fully Convolutional Network for segmentation

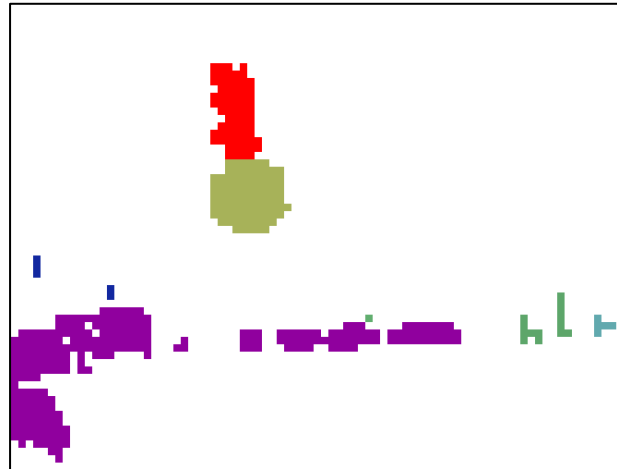


# Up-Scaling-Layer

encoded input

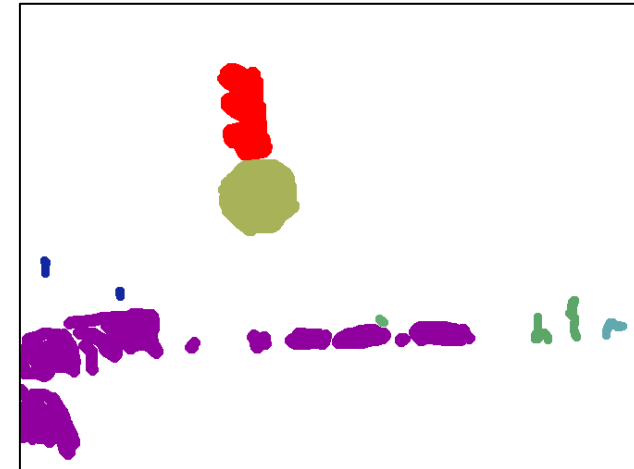


upscaled



$h \times w \times c$

Ground Truth



$h \times w \times c$

$$\frac{h}{16} \times \frac{w}{16} \times c$$



# FCN-Upscaling: Example

## Pseudo Code

Conv 40, 3x3, padding=same

ReLU

MaxPool 2x2

Conv 40, 3x3, padding=same

ReLU

MaxPool 2x2

Conv 80, 3x3, padding=same

ReLU

MaxPool 2x2

UpSampling 8x8

Conv2D n\_classes, 1x1

Softmax

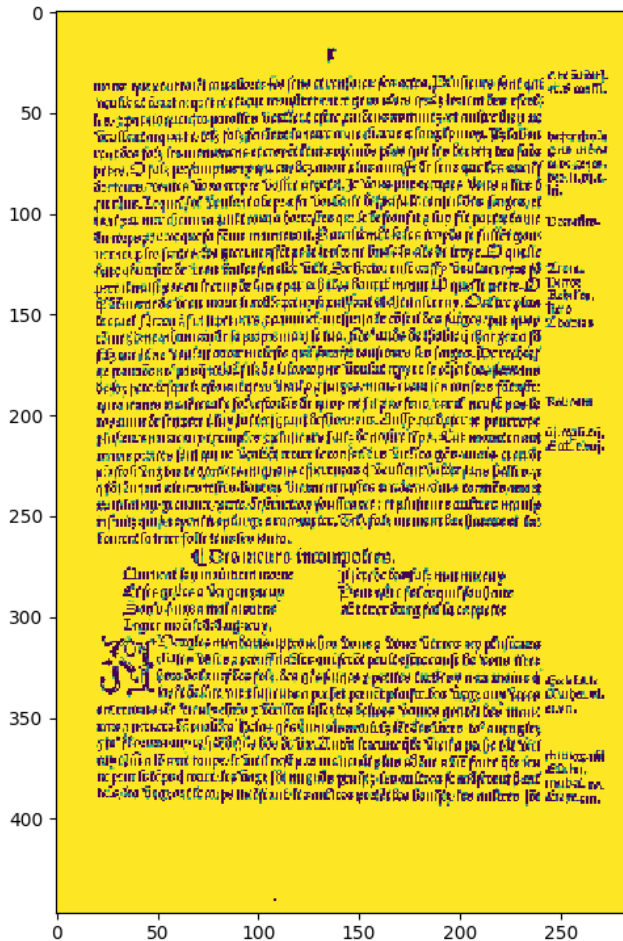
## Keras

```
model = Sequential([
    Conv2D(40, (3, 3), padding='same', activation='relu',
input_shape=(None, None, 1)),
    MaxPooling2D(pool_size=(2, 2), strides=2),
    Conv2D(40, (4, 4), padding='same', activation='relu'),
    MaxPooling2D(pool_size=(2, 2), strides=2),
    Conv2D(80, (4, 4), padding='same', activation='relu'),
    MaxPooling2D(pool_size=(2, 2), strides=2),
    UpSampling2D((8, 8), interpolation='nearest'),
    Conv2D(num_classes, (1, 1)),
    Activation('softmax'),
])
```

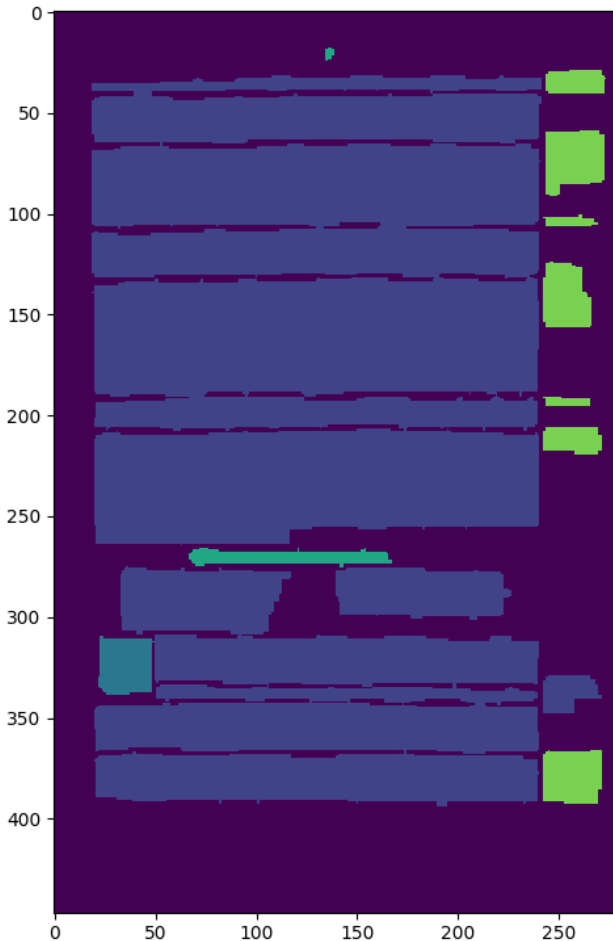


# FCN-Upscaling: Example

Input (binary image)



Ground Truth (Label image)

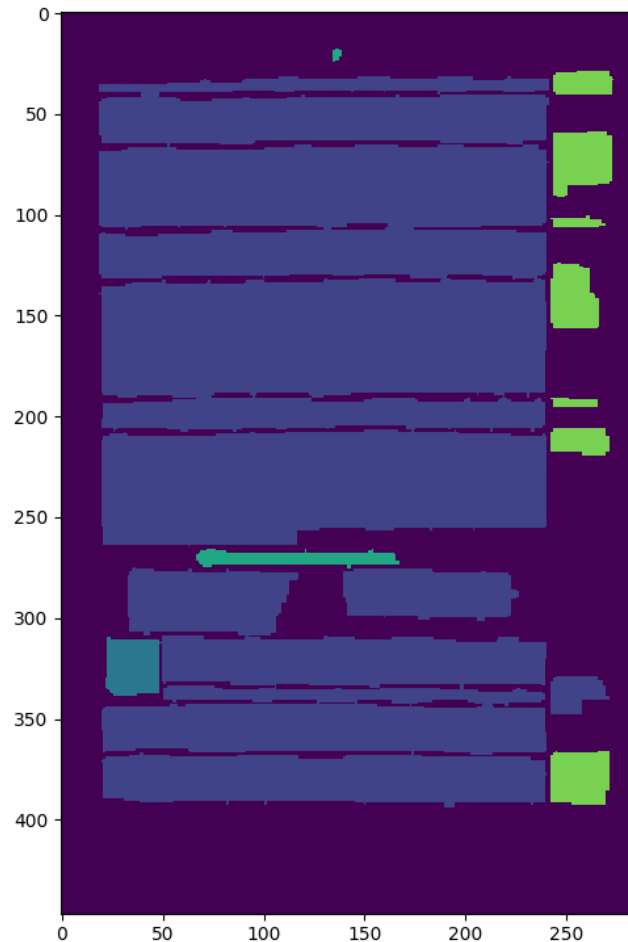


Prediction

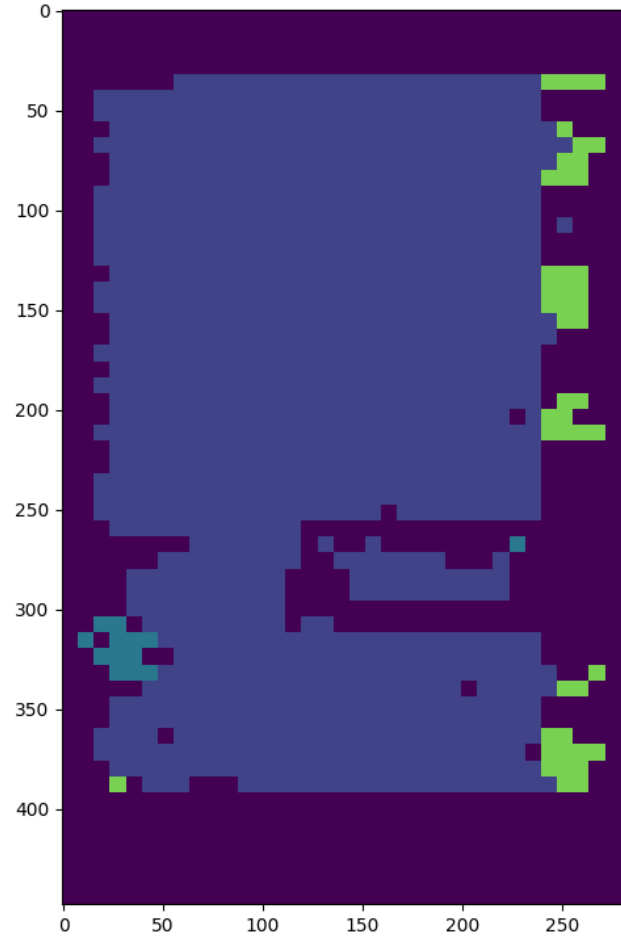


# FCN-Upscaling: Example

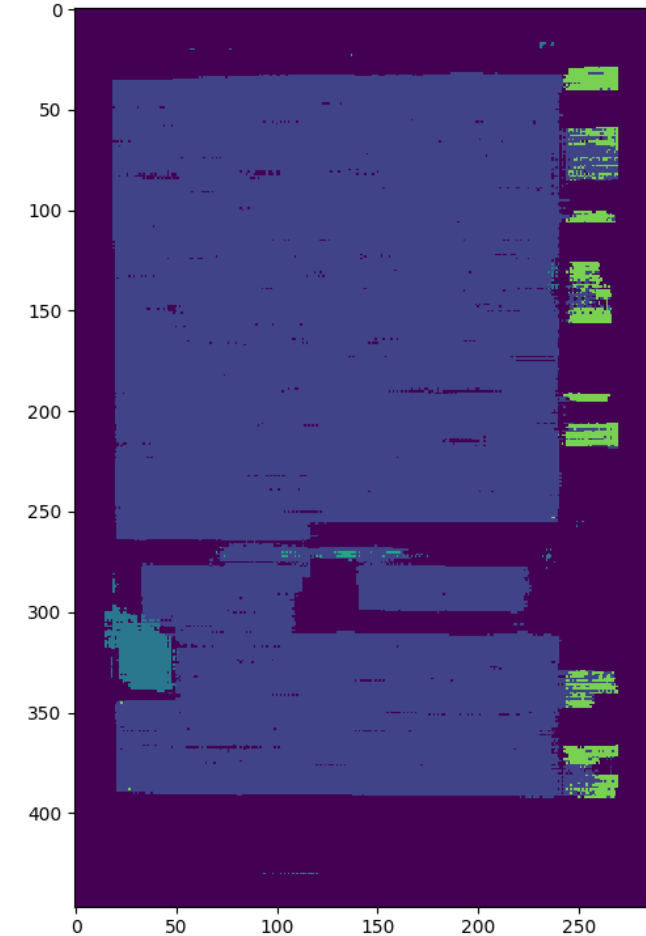
Ground Truth (Label image)



Prediction FCN (0.9s)



Prediction Sliding Window (6s)





# Up-Scaling-Layer

## Pros

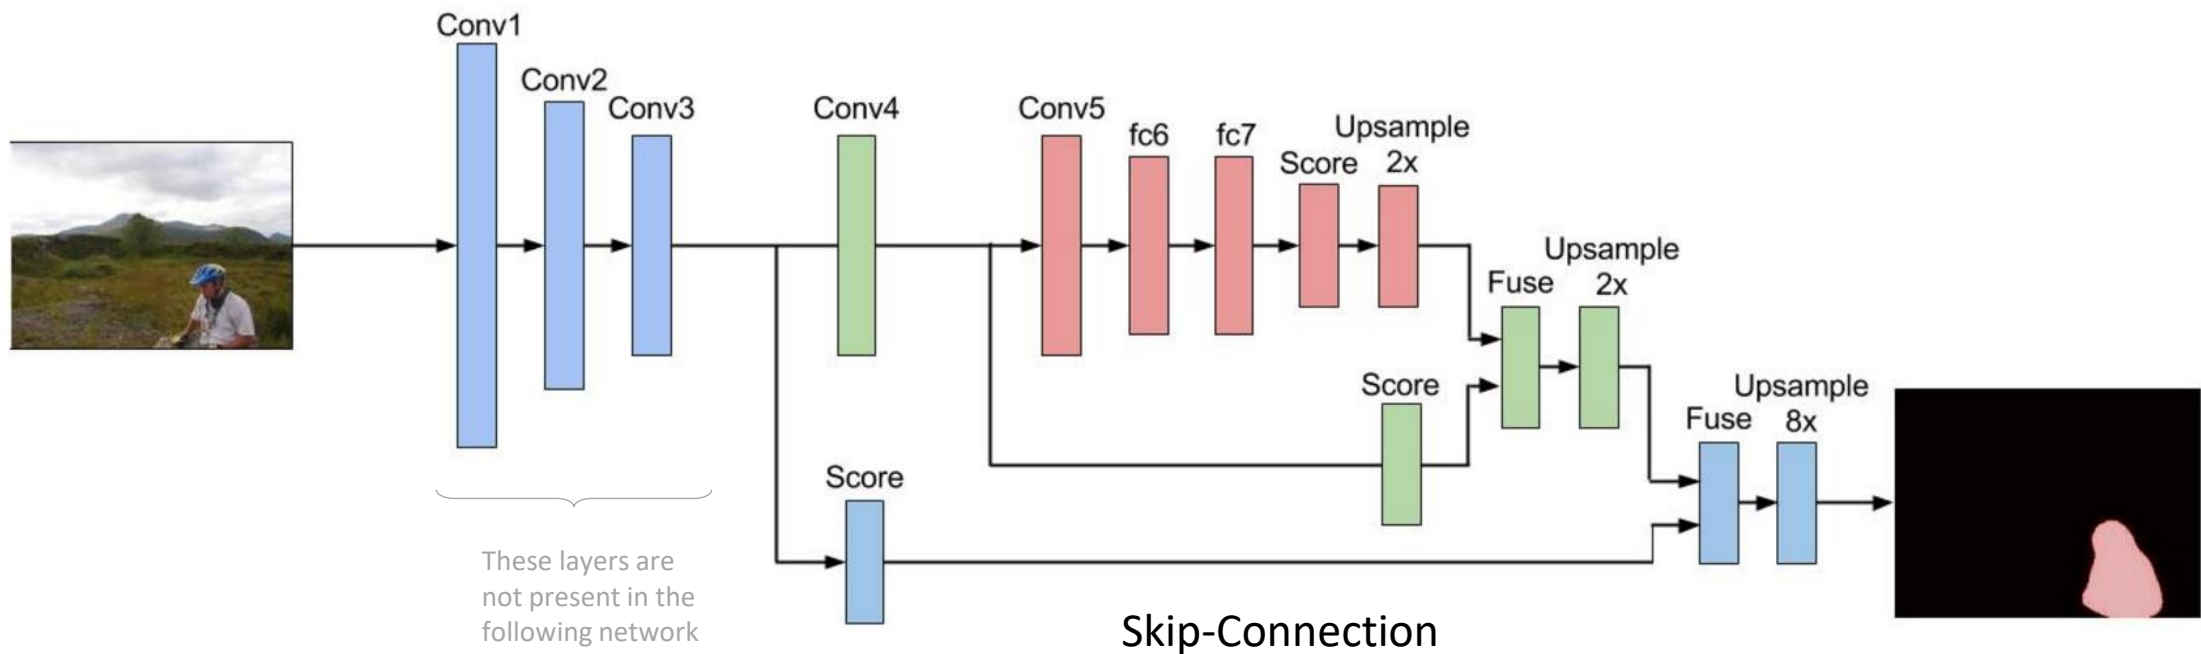
- Easy to implement (error is summed in the backward pass)
- No parameters  $\Rightarrow$  extremely fast

## Cons

- Small objects absorbed
- Only very rough regions



# FCN: Skip-Connections



# Keras: Skip-Connections

```
input = Input((None, None, 1))

conv1 = Conv2D(40, (3, 3), padding='same',
activation='relu')(input)

pool1 = MaxPooling2D(pool_size=(2, 2),
strides=2)(conv1)

conv2 = Conv2D(80, (4, 4), padding='same',
activation='relu')(pool1)

pool2 = MaxPooling2D(pool_size=(2, 2),
strides=2)(conv2)

conv3 = Conv2D(80, (4, 4), padding='same',
activation='relu')(pool2)

pool3 = MaxPooling2D(pool_size=(2, 2),
strides=2)(conv3)

fc4 = Conv2D(160, (1, 1), padding='same',
activation='relu')(pool3)
```

```
fc5 = Conv2D(160, (1, 1), padding='same',
activation='relu')(fc4)

up1 = UpSampling2D((2, 2), fc5)

fuse2 = Concatenate(axis=-1)([up1, conv3])

up2 = UpSampling2D((2, 2), fuse2)

fuse3 = Concatenate(axis=-1)([up2, conv2])

up3 = UpSampling2D((2, 2), fuse3)

fuse4 = Concatenate(axis=-1)([up3, conv1])

logits = Conv2D(num_classes, (1, 1))(fuse4)

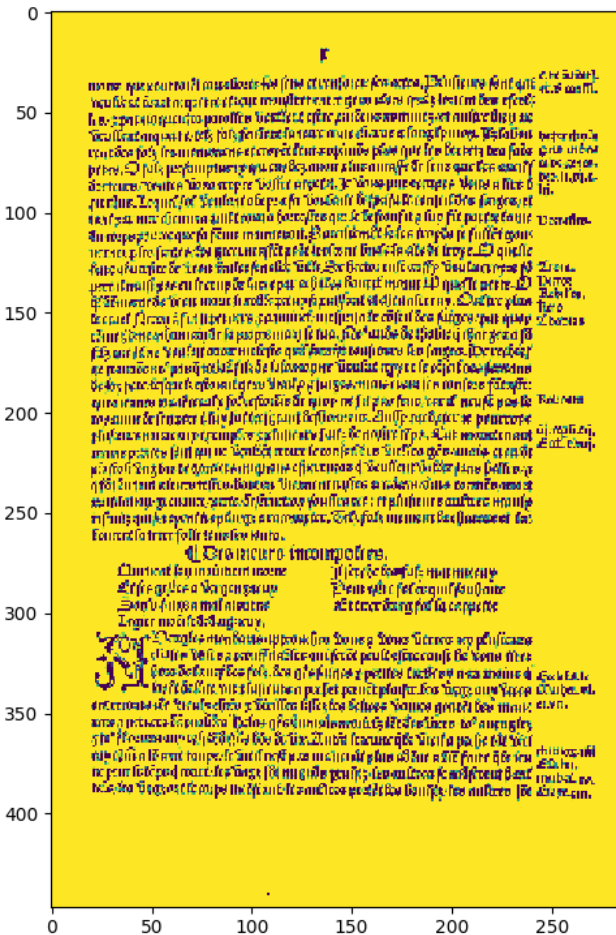
softmax = Activation('softmax')(logits)

model = Model(inputs=input, outputs=softmax)
```

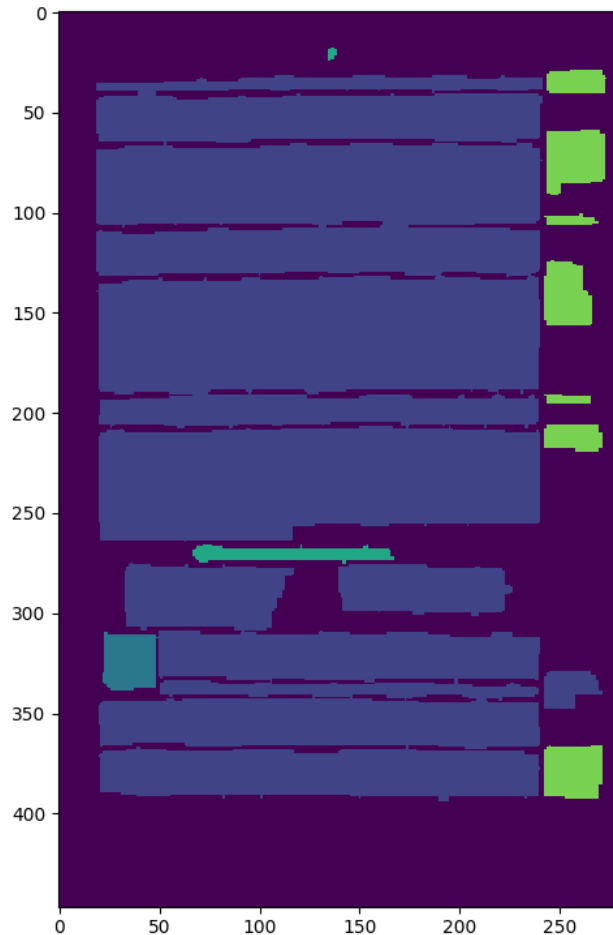


# FCN-Upscaling-Skip-Connections: Example

Input (Binary Image)



Ground Truth (Label image)



Prediction

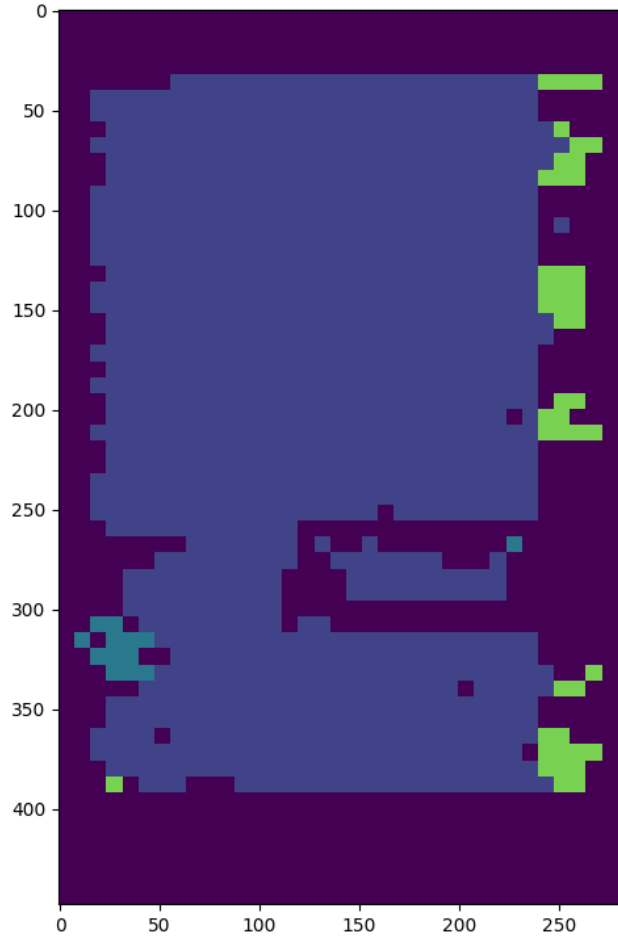


# FCN-Upscaling-Skip-Connections: Example

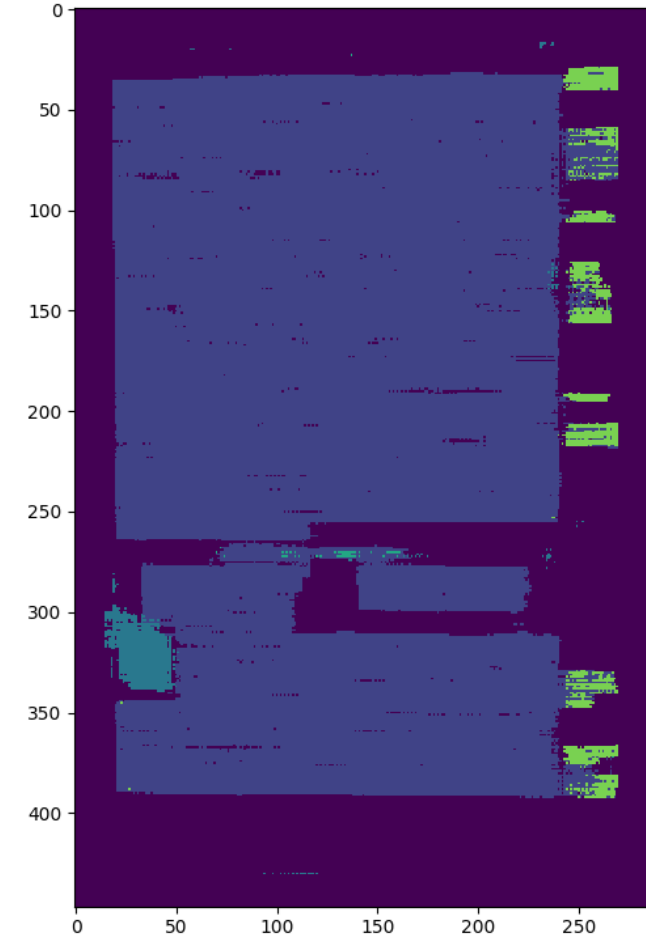
Prediction Skip-FCN (0.91s)



Prediction FCN (0.86s)

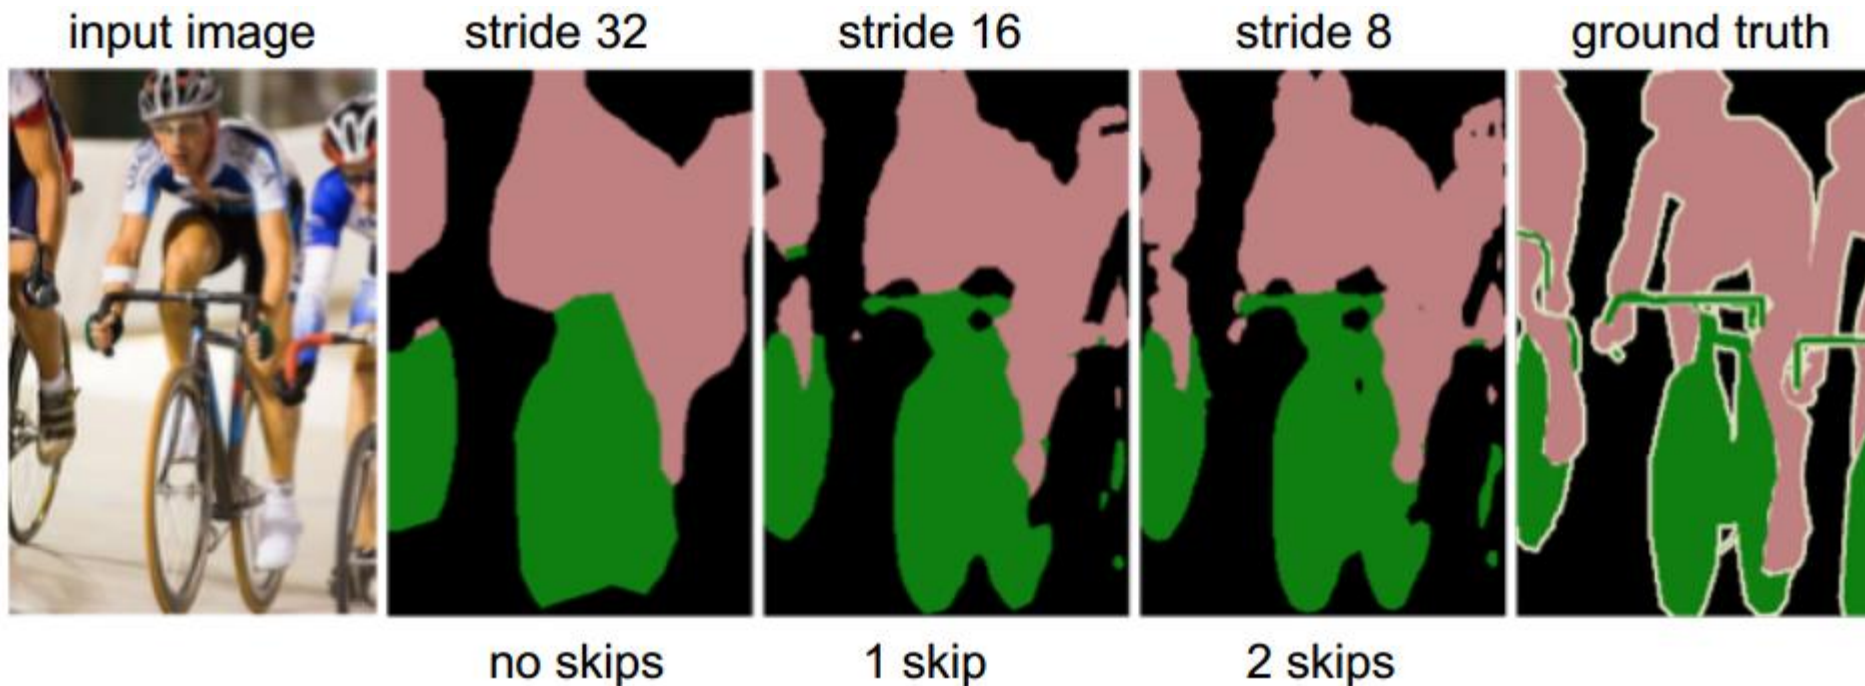


Prediction Sliding Window (6s)



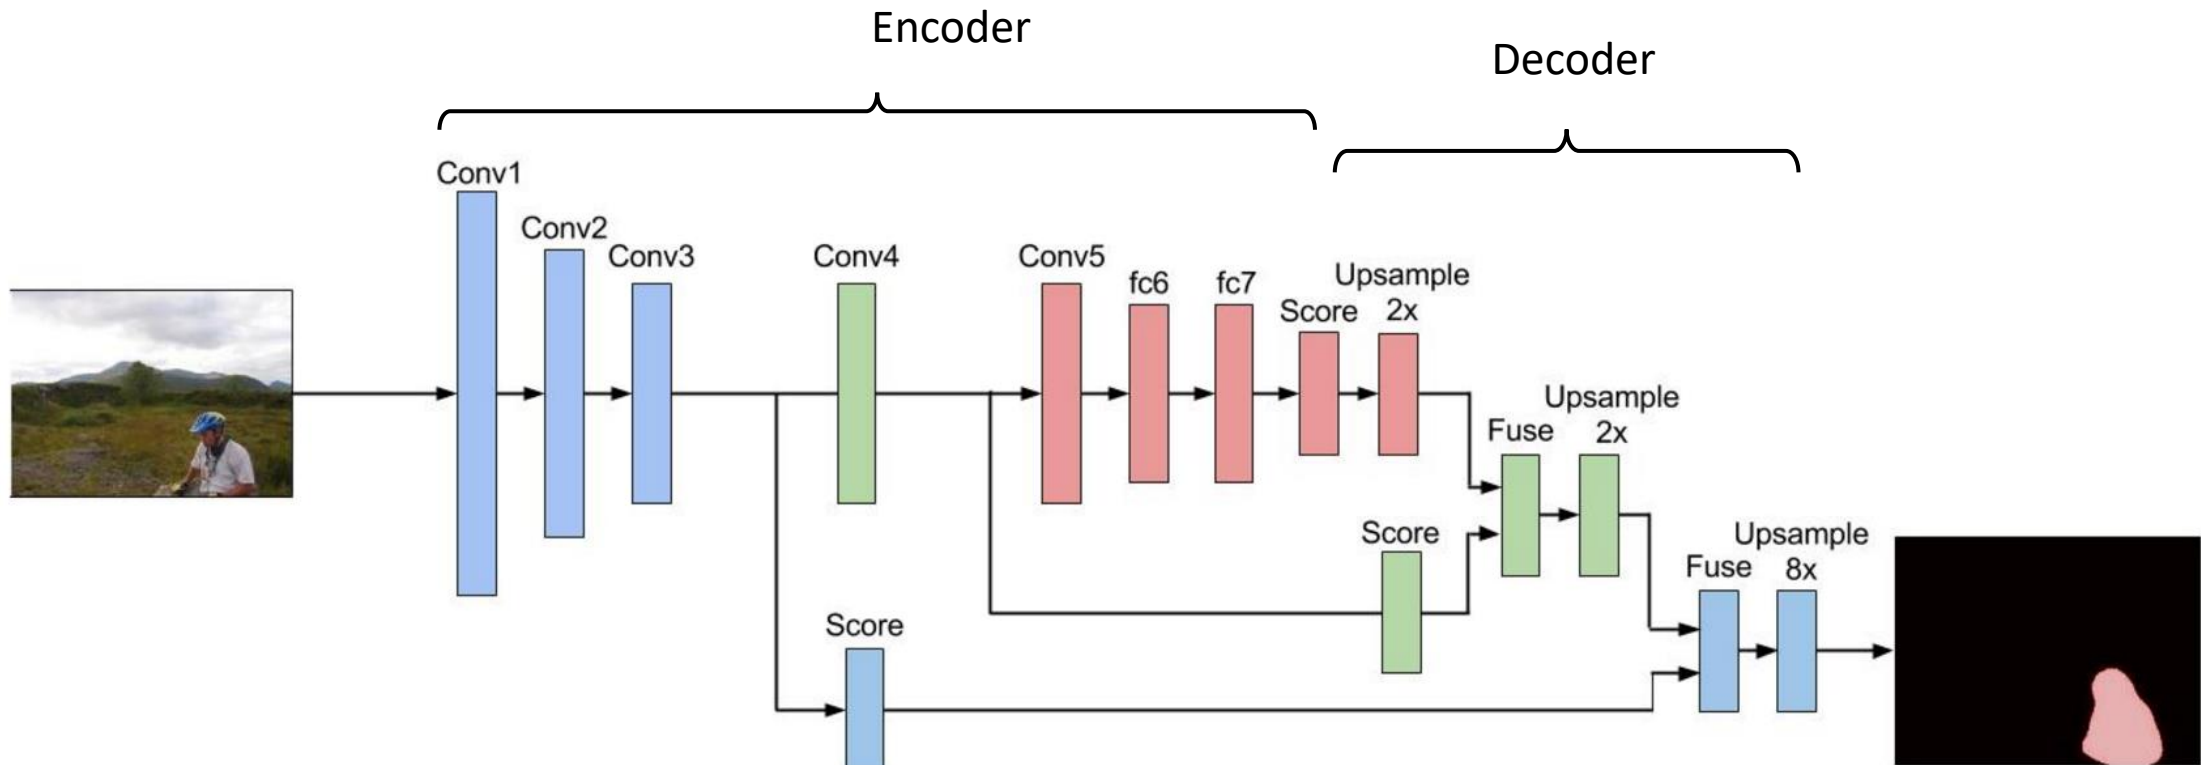
# FCN-Upscaling-Skip-Connections:

- Skip connections allow resolving finer structures
- The complete Encoder-Decoder generates rough structures





# FCN: Encoder/Decoder naming



# Insertion: FCN for compression

**Input** image with high resolution with full information (e.g. bmp)



Encoder  
algorithm

JPG

Decoder  
algorithmus

**Decoded** image as close as possible to the original



Compressed  
representation of  
the image

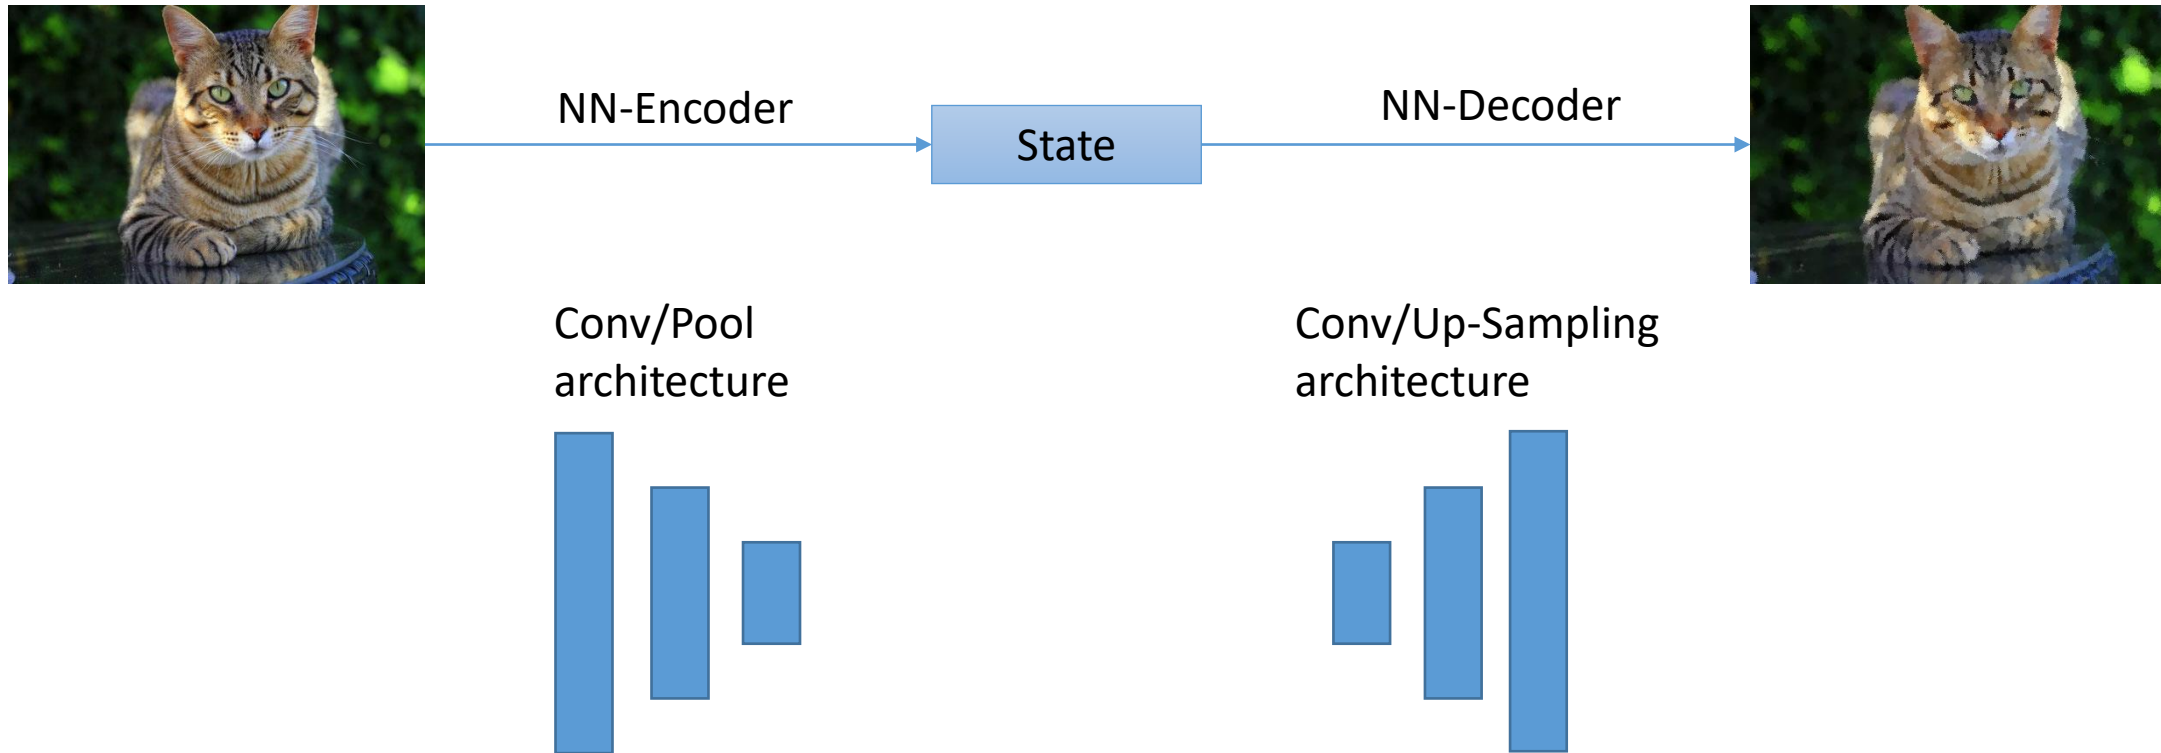
NN-Encoder

State

NN-Decoder



# Insertion: FCN for compression



# Up-Sampling Alternatives

Transposed Convolution, Unpooling



# Up-Sampling: Methods

- Up-Scaling (no parameters)
- Transposed convolution (sometimes also deconvolution)
- Un-Pooling („Inverse pooling“)



# Transposed Convolution

Example U-Net

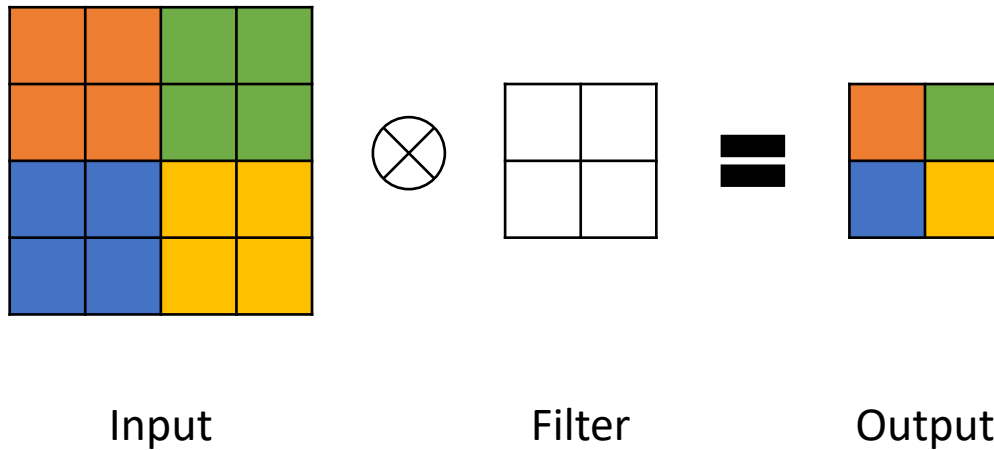




# Convolution

Example: Convolution

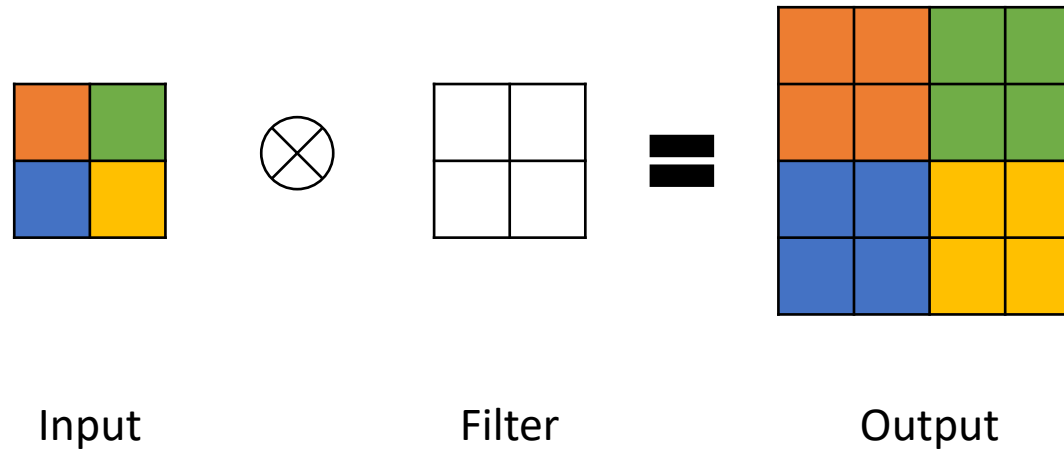
- Kernel-Size  $2 \times 2 \times k$
- Stride 2



# Transposed Convolution

Example: Transposed Convolution

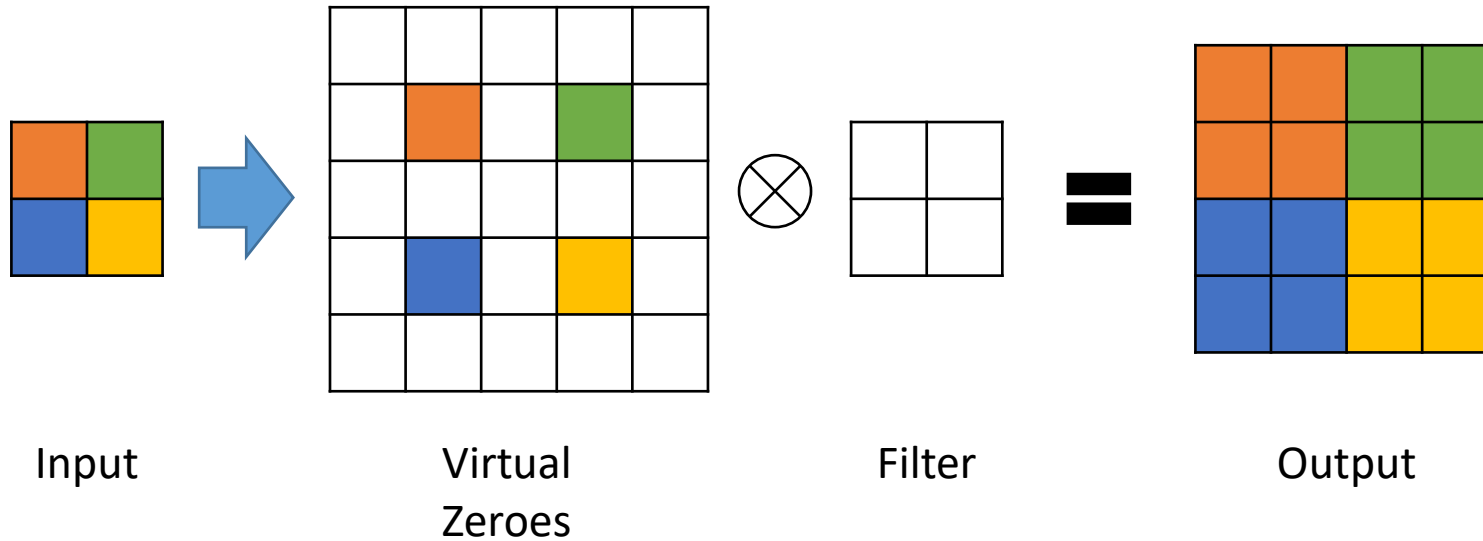
- Kernel-Size  $2 \times 2 \times k$
- Stride 2



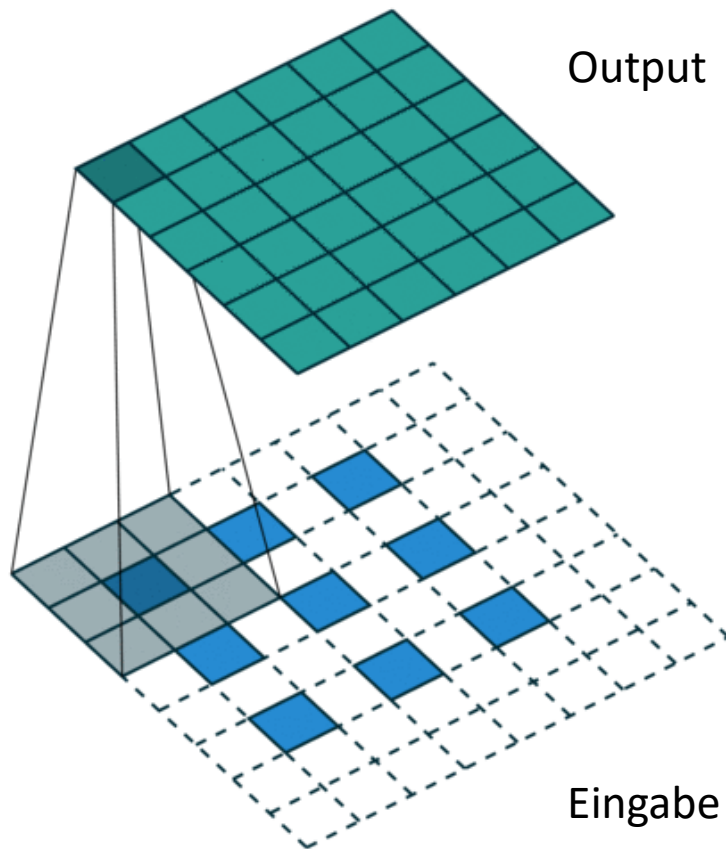
# Transposed Convolution

Example: Transposed Convolution

- Kernel-Size  $2 \times 2 \times k$
- Stride 2



# Transposed Convolution



- Transposed convolution here with kernel size 3x3, stride 2 and padding generates an output image with doubled dimensions of the input image
- Animation should seem familiar



# Comparing Convolutions

## Convolution

- Padding
  - Full padding
  - No Padding
  - Half Padding
- Stride  $s$ : Resolution decreased by factor  $s$
- Conv in Forward-Pass becomes zu transposed Conv in Backward-Pass

## Transposed Convolution

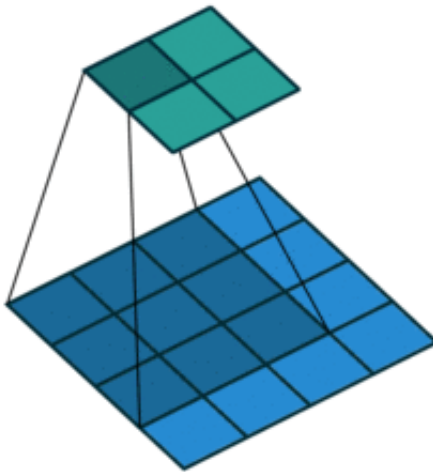
- Padding
  - No Padding
  - Full Padding
  - Half Padding
- Stride  $s$ : Resolution increased by factor  $s$
- Transposed Conv in Forward-Pass becomes Conv in Backward-Pass



# Tranposed Convolution Examples: Standard

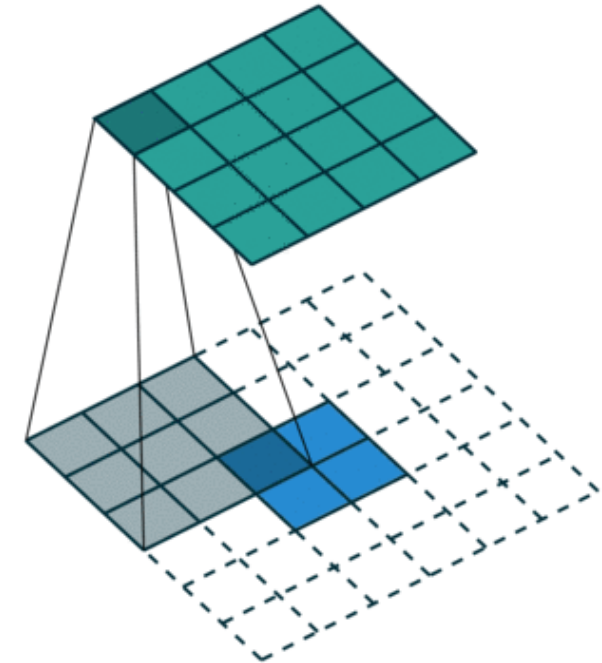
## Normal Convolution

- 3x3 Filter, Stride 1
- 4x4 Input
- 2x2 Output



## Transposed Convolution

- 3x3 Filter, Stride 1
- 2x2 Input
- 4x4 Output

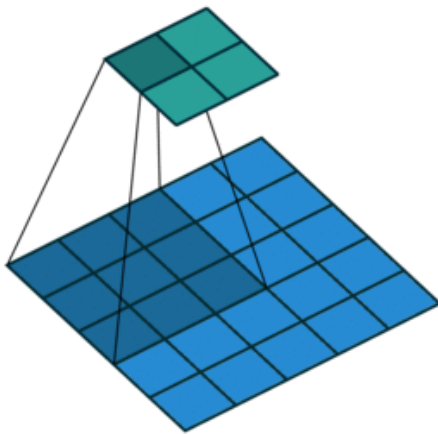




# Tranposed Convolution Examples: Stride

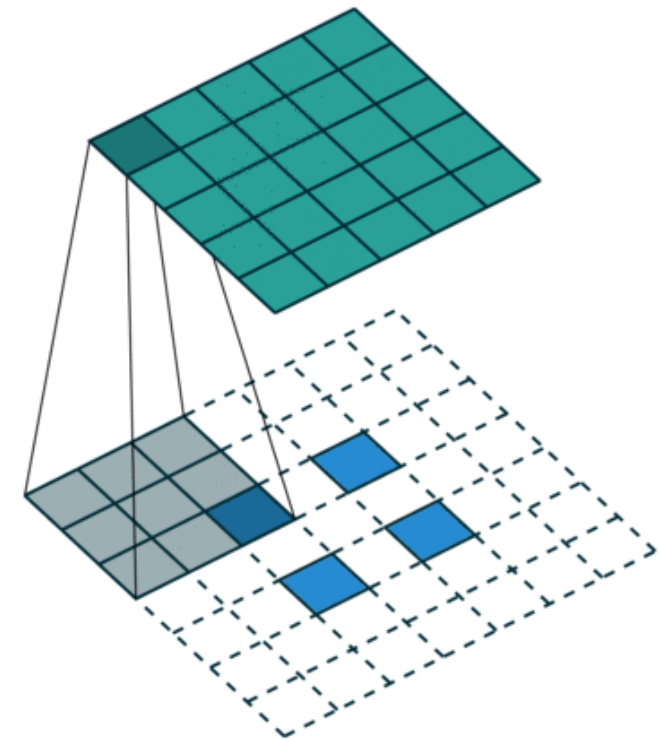
## Normal Convolution

- 3x3 Filter, Stride 2
- 5x5 Input
- 2x2 Output



## Transposed Convolution

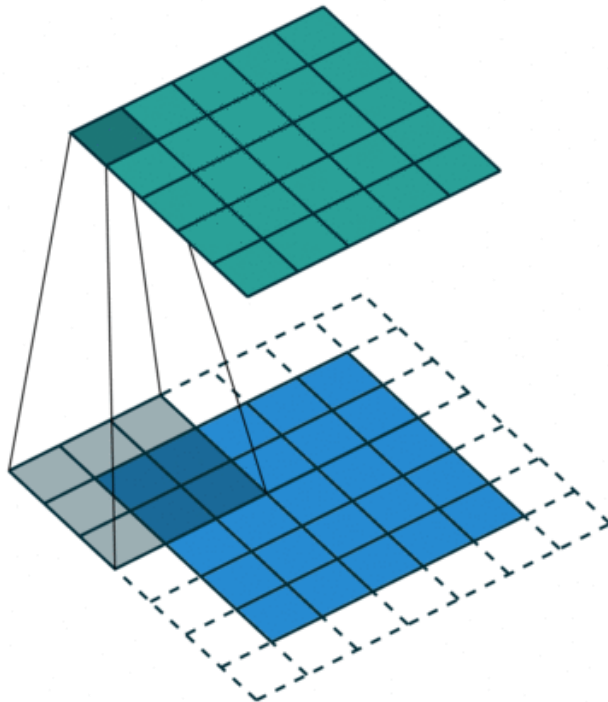
- 3x3 Filter, Stride 2
- 2x2 Input
- 5x5 Output



# Tranposed Convolution Examples: Padding

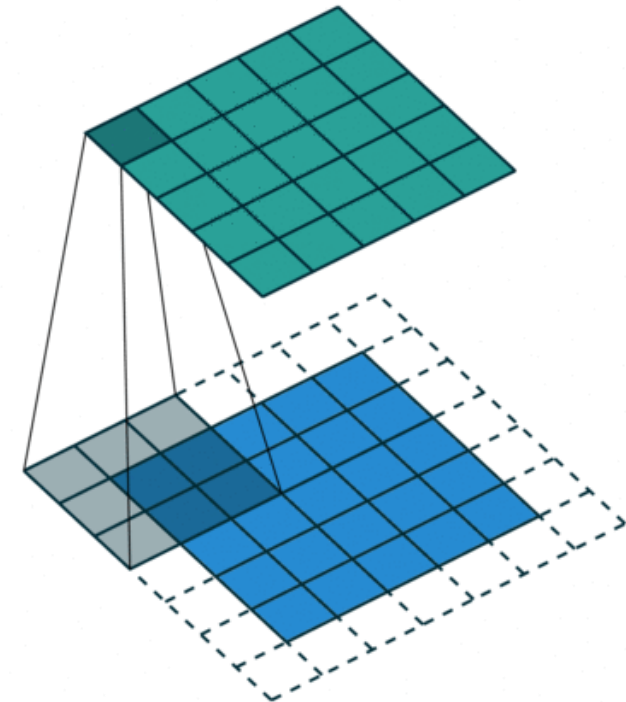
## Normal Convolution

- 3x3 Filter, Stride 1
- 5x5 Input
- 5x5 Output



## Transposed Convolution

- 3x3 Filter, Stride 1
- 5x5 Input
- 5x5 Output



# FCN-Conv<sup>T</sup>-Skip-Connections: Architecture

```
input = Input((None, None, 1))

conv1 = Conv2D(40, (3, 3), padding='same',
activation='relu')(input)

pool1 = MaxPooling2D(pool_size=(2, 2),
strides=2)(conv1)

conv2 = Conv2D(80, (4, 4), padding='same',
activation='relu')(pool1)

pool2 = MaxPooling2D(pool_size=(2, 2),
strides=2)(conv2)

conv3 = Conv2D(80, (4, 4), padding='same',
activation='relu')(pool2)

pool3 = MaxPooling2D(pool_size=(2, 2),
strides=2)(conv3)

fc4 = Conv2D(160, (1, 1), padding='same',
activation='relu')(pool3)
```

```
fc5 = Conv2D(160, (1, 1), padding='same',
activation='relu')(fc4)

up1 = Conv2DTransposed(80, (3, 3), 2,
padding='same')(fc5)

fuse2 = Concatenate(axis=-1)([up1, conv3])

up2 = Conv2DTransposed(80, (3, 3), 2,
padding='same')(fuse2)

fuse3 = Concatenate(axis=-1)([up2, conv2])

up3 = Conv2DTransposed(40, (3, 3), 2,
padding='same')(fuse3)

fuse4 = Concatenate(axis=-1)([up3, conv1])

logits = Conv2D(num_classes, (1, 1))(fuse4)

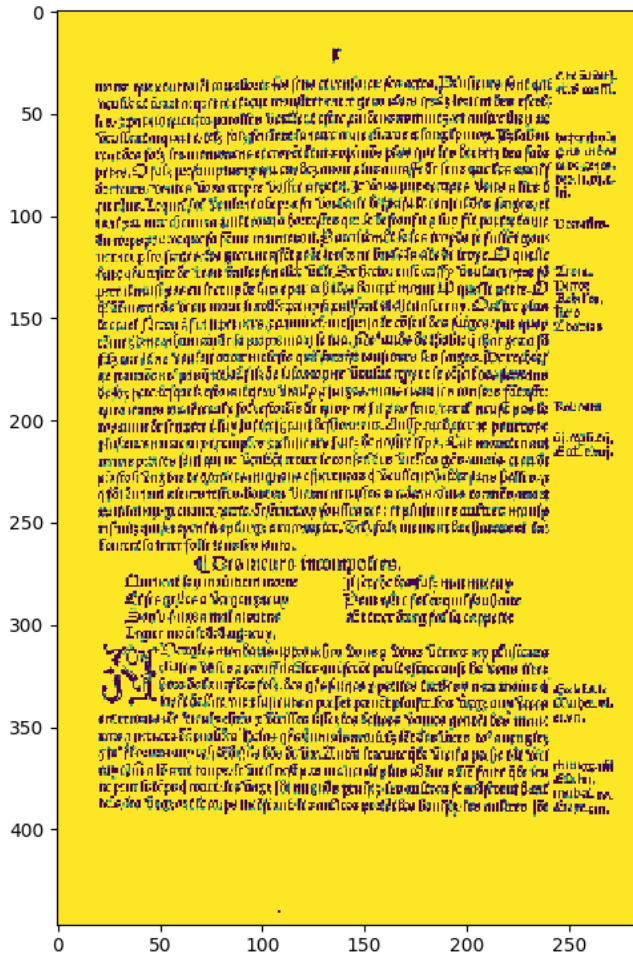
softmax = Activation('softmax')(logits)

model = Model(inputs=input, outputs=softmax)
```

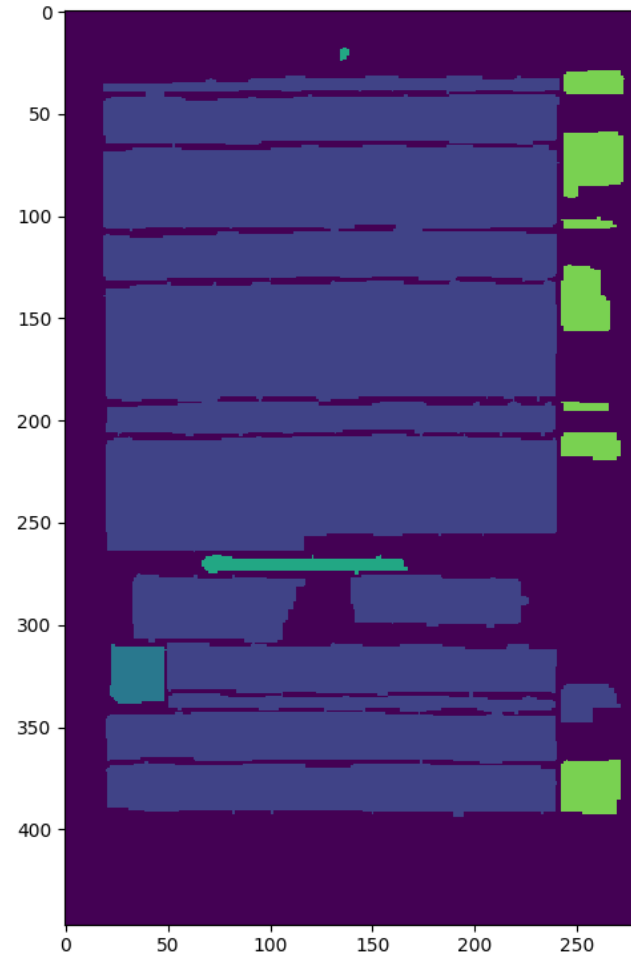


# FCN-Conv<sup>T</sup>-Skip-Connections: Example

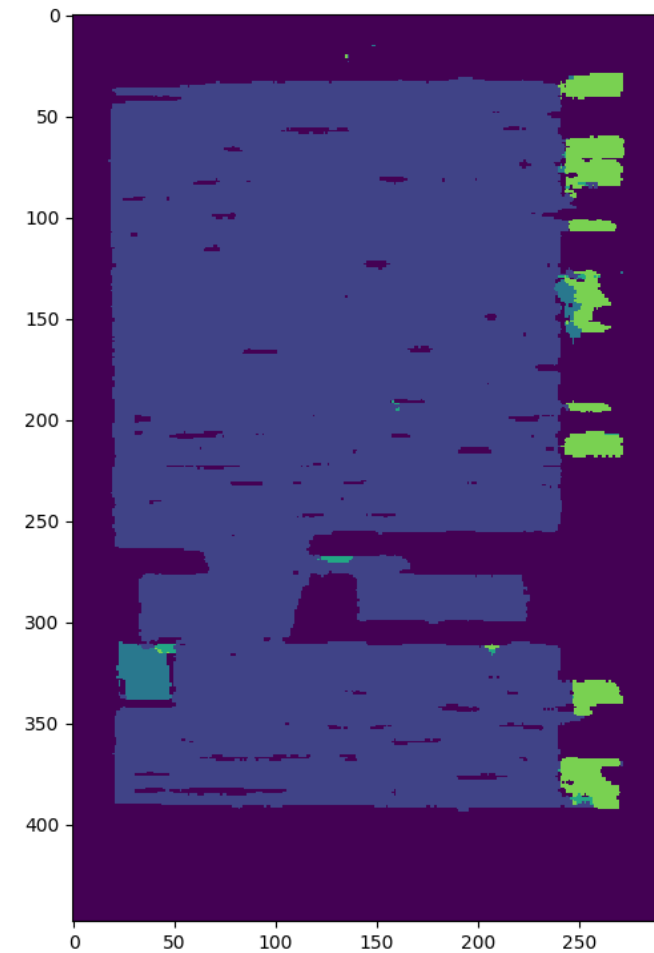
Input (Binary Image)



Ground Truth (Label Image)



Prediction (1.1s)

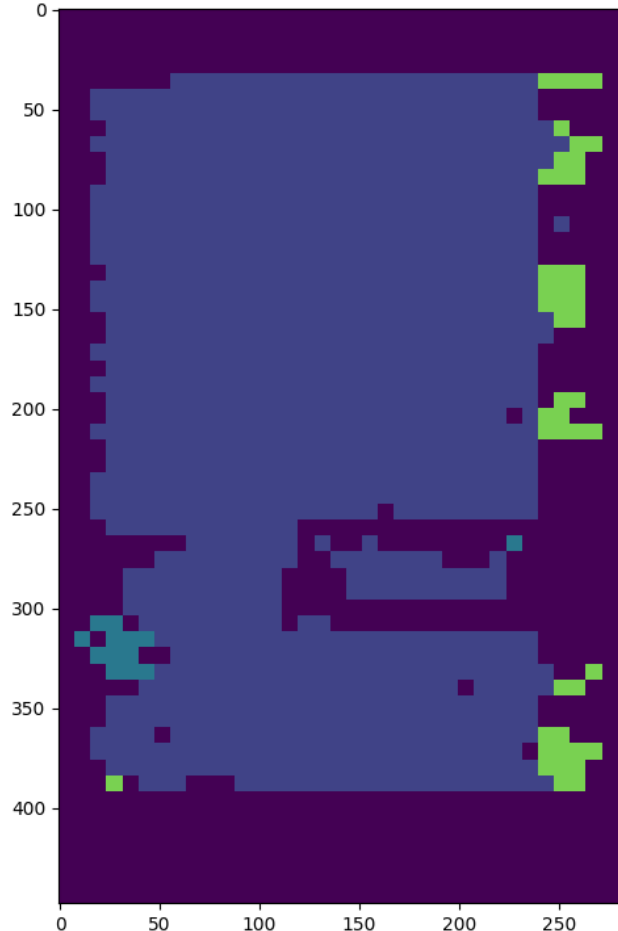


# FCN-Conv<sup>T</sup>-Skip-Connections: Example

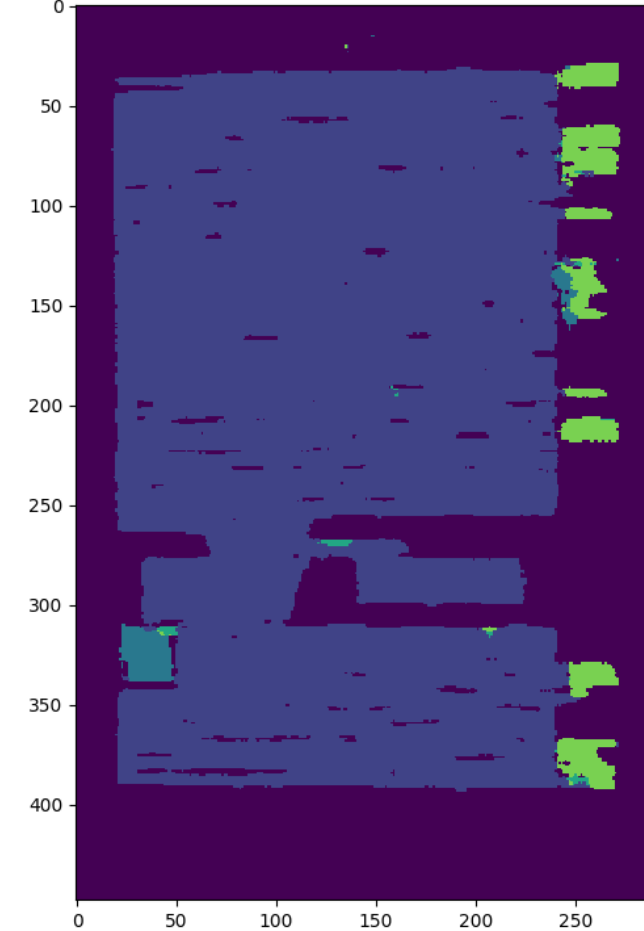
Prediction Skip-FCN (0.91s)



Prediction FCN (0.86s)



Prediction Transposed Conv (1.1s)



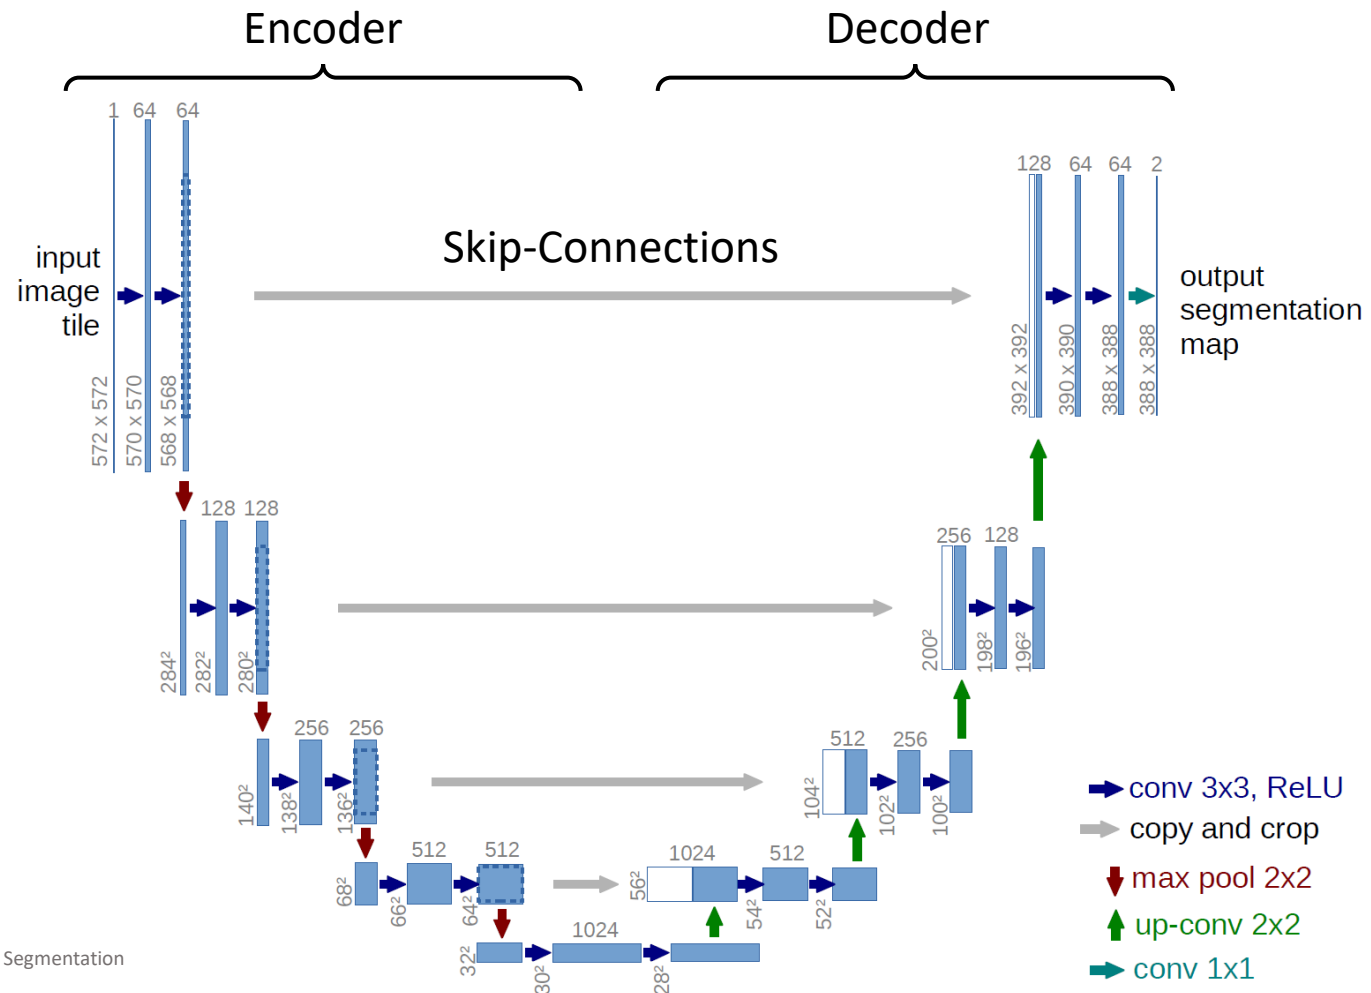
# U-Net

- So far: No additional conv layers after up-sampling (except for output layer)
- In reality (actual networks): Decoder consists of several more conv layers, analogous to the encoder
- Popular network: U-Net





# U-Net



U-Net: Convolutional Networks for Biomedical Image Segmentation  
Olaf Ronneberger, Philipp Fischer, Thomas Brox



# Up-Sampling Alternatives

Unpooling

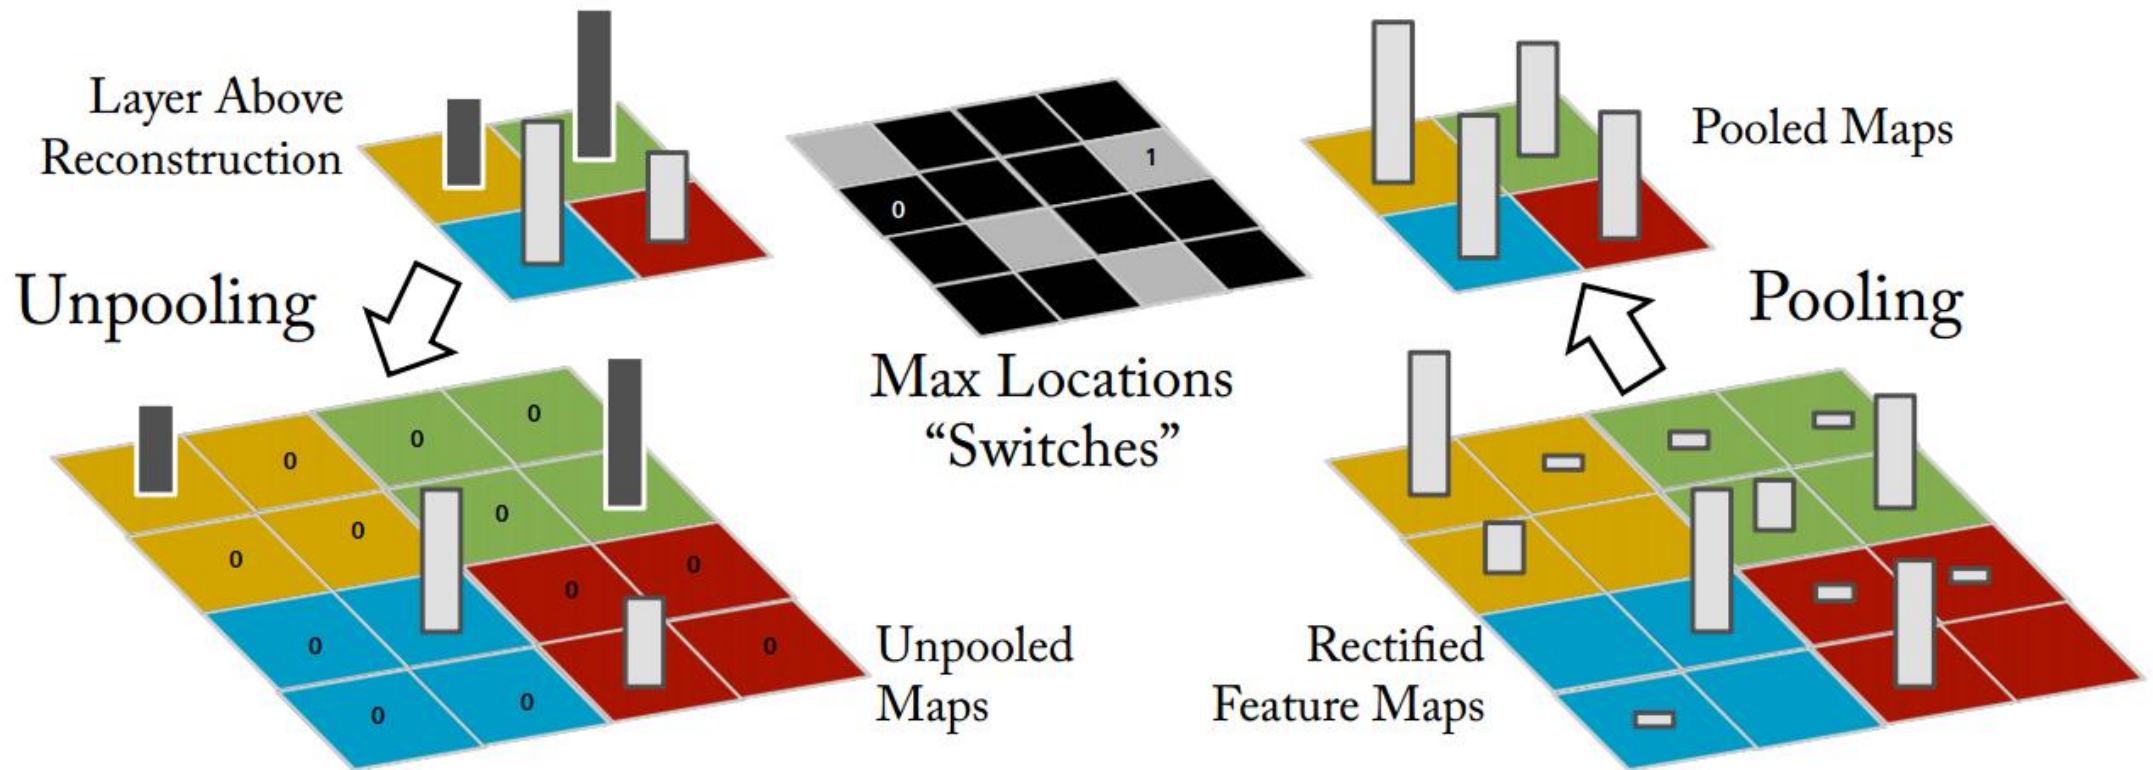


# Unpooling

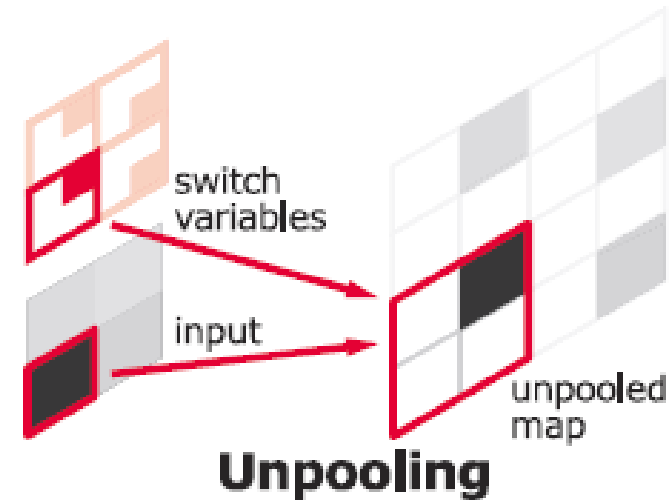
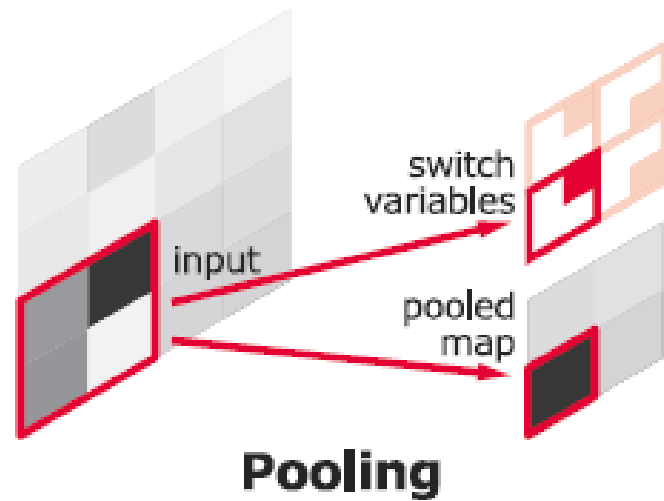
- Unpooling as alternative to Transposed Convolution layers
- A pooling layer is required for every unpooling layer
- Max pooling is non reversible, but analogous to backward pass:
  - Save, where the weight came from
  - Forward the input value to this position



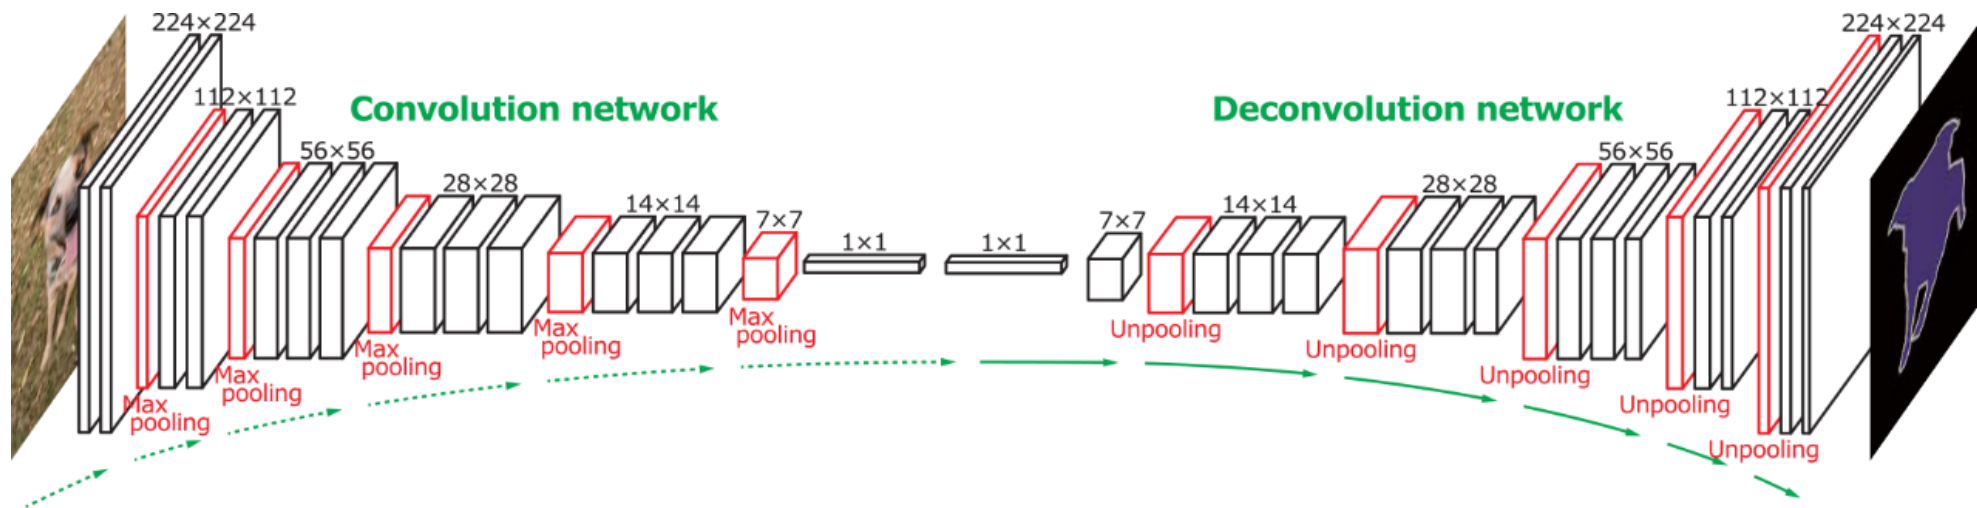
# Unpooling



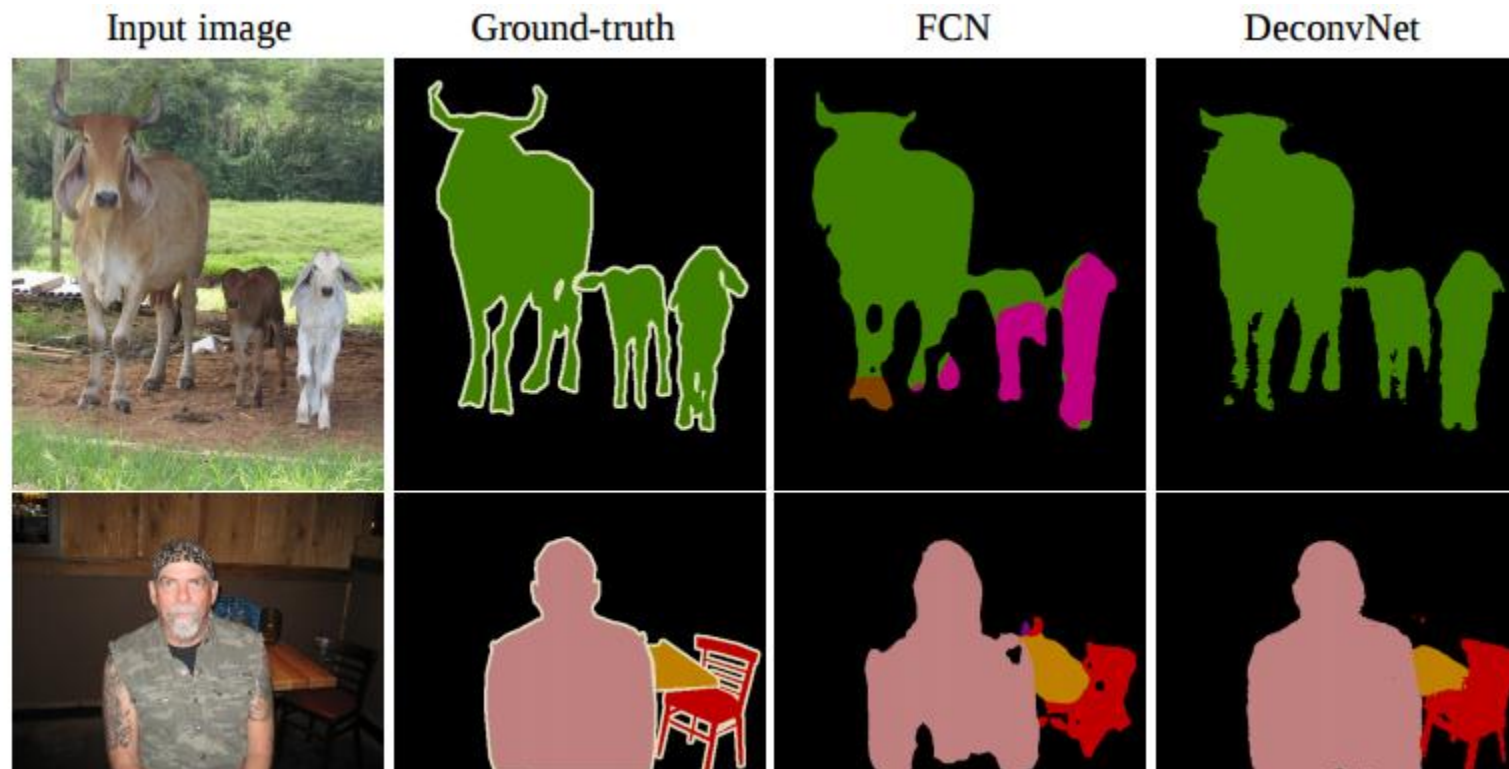
# Unpooling



# Example: Deconvnet



# Example: Deconvnet





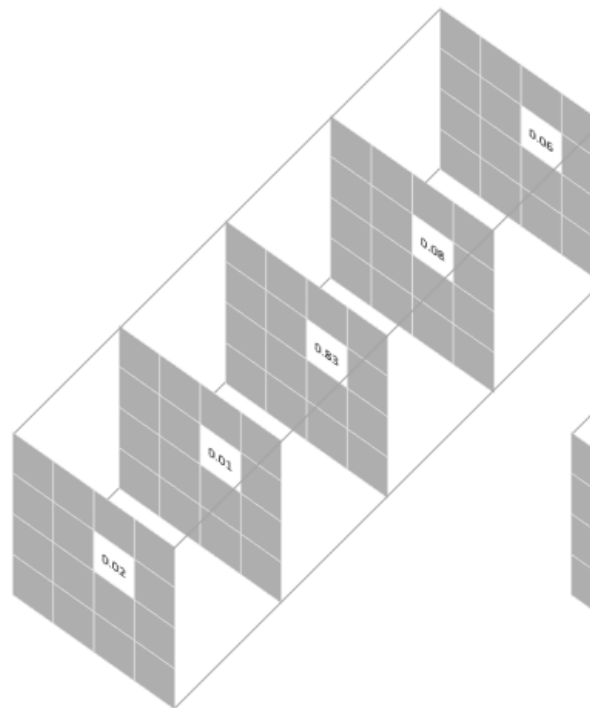
# Training an FCN

Loss, Evaluation and Transfer-Learning

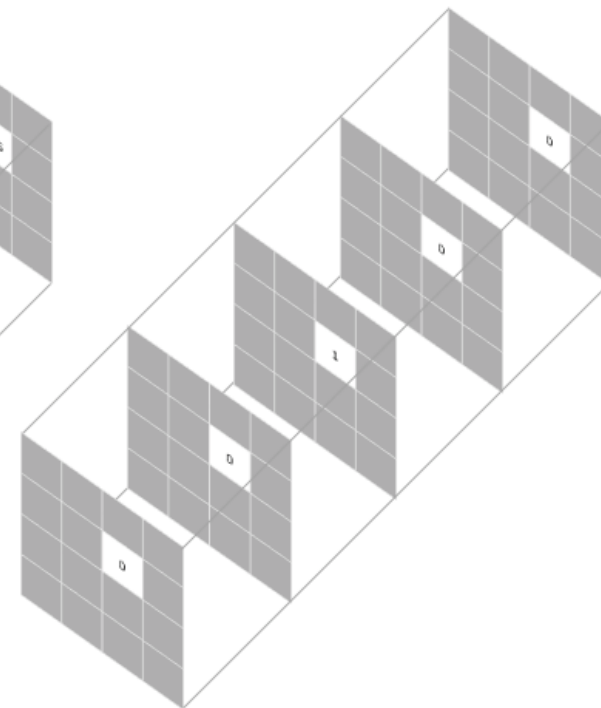


# LOSS

- Same loss function as with classification possible
- Difference: Sum and average across **all pixels**
- Problem: Some classes overrepresented (background)



Prediction for a selected pixel



Target for the corresponding pixel

Pixel-wise loss is calculated as the log loss, summed over all possible classes

$$-\sum_{\text{classes}} y_{\text{true}} \log(y_{\text{pred}})$$

This scoring is repeated over all **pixels** and averaged

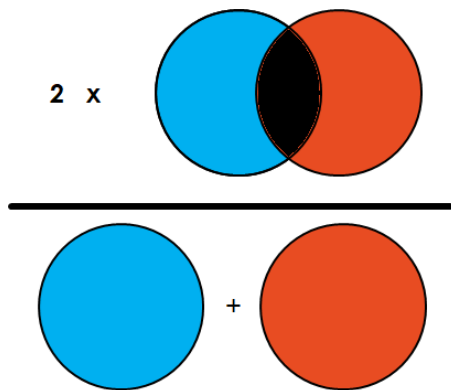


# Evaluation metrics

## Dice-Koeffizient

- Effektiv F1-Score

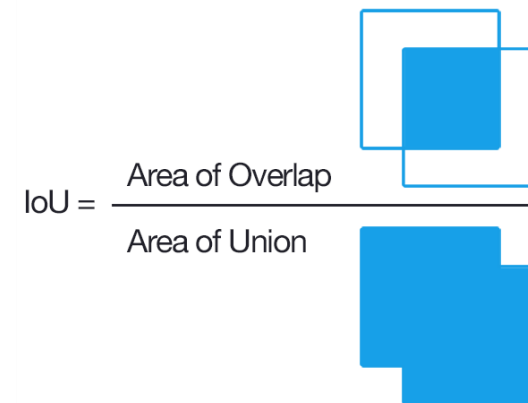
$$\frac{2 \cdot |A \cap B|}{|A| + |B|} = \frac{2 \cdot TP}{2 \cdot TP + FP + FN}$$



## Intersection over Union

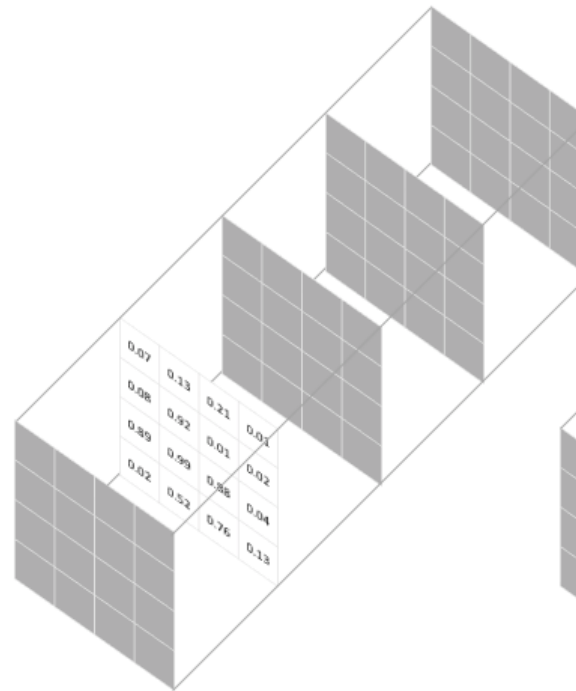
- Also known as Jaccard Index

$$\frac{|A \cap B|}{|A| + |B| - |A \cap B|} = \frac{TP}{TP + FP + FN}$$

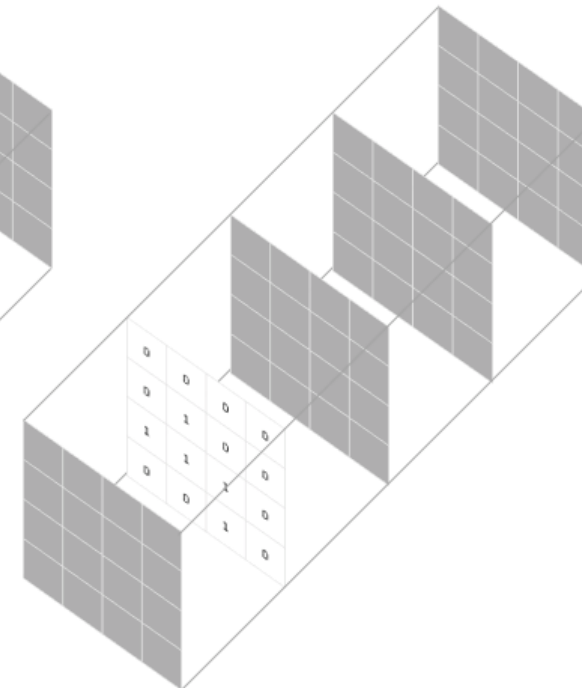


# Variation: Dice-Loss

- Goal: Maximizing Dice-Score
- Formulate as differentiable loss function
- $y_{\text{true}}$  are GT labels
- $y_{\text{pred}}$  are predicted probabilities
- Computation happens per prediction mask
- Better class balance



Prediction for a selected class



Target for the corresponding class

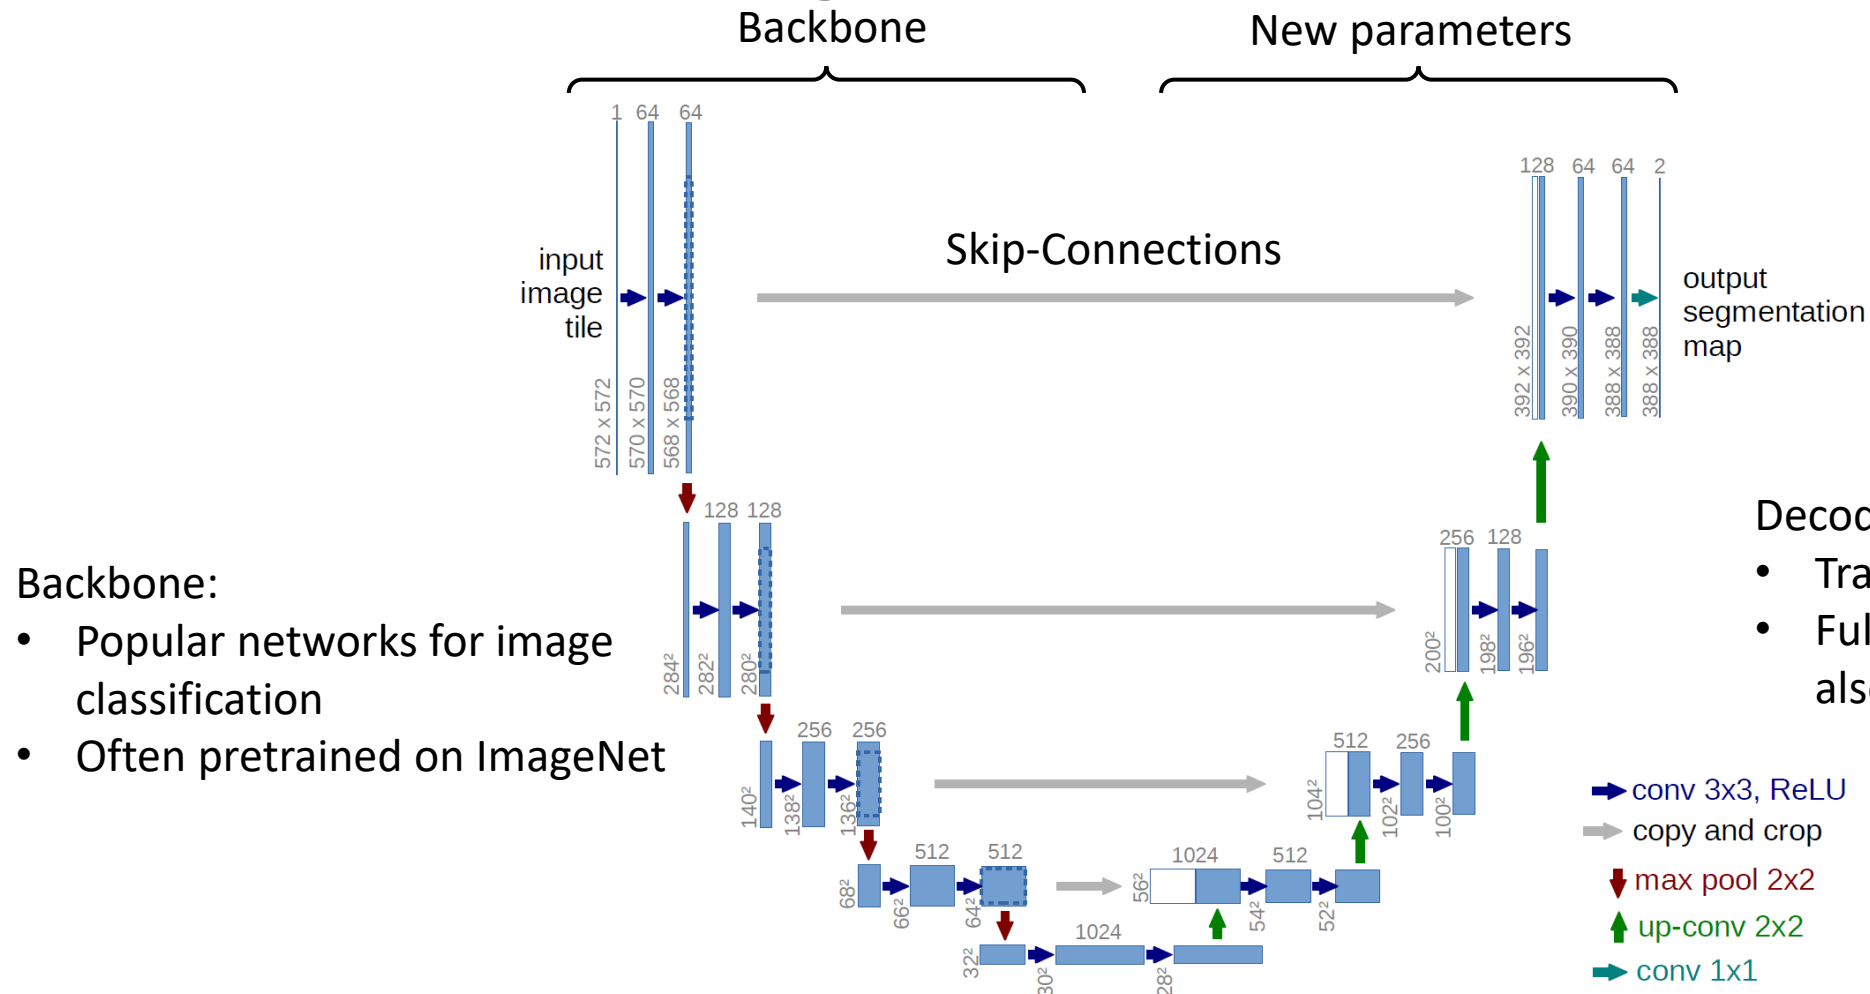
Soft Dice coefficient is calculated for each class mask

$$1 - \frac{2 \sum_{\text{pixels}} y_{\text{true}} y_{\text{pred}}}{\sum_{\text{pixels}} y_{\text{true}}^2 + \sum_{\text{pixels}} y_{\text{pred}}^2}$$

This scoring is repeated over all **classes** and averaged



# Transfer Learning



Backbone:

- Popular networks for image classification
- Often pretrained on ImageNet

Decoder parameters:

- Training from Scratch
- Fully pretrained networks also available



# Transfer-Learning: Datasets

- [COCO 2018 Stuff Segmentation](#)
- [BDD100K: A Large-scale Diverse Driving Video Database](#)
- [Cambridge-driving Labeled Video Database \(CamVid\)](#)
- [Cityscapes Dataset](#)
- [Mapillary Vistas Dataset](#)
- [Apolloscape Scene Parsing](#)



# Instance Segmentation

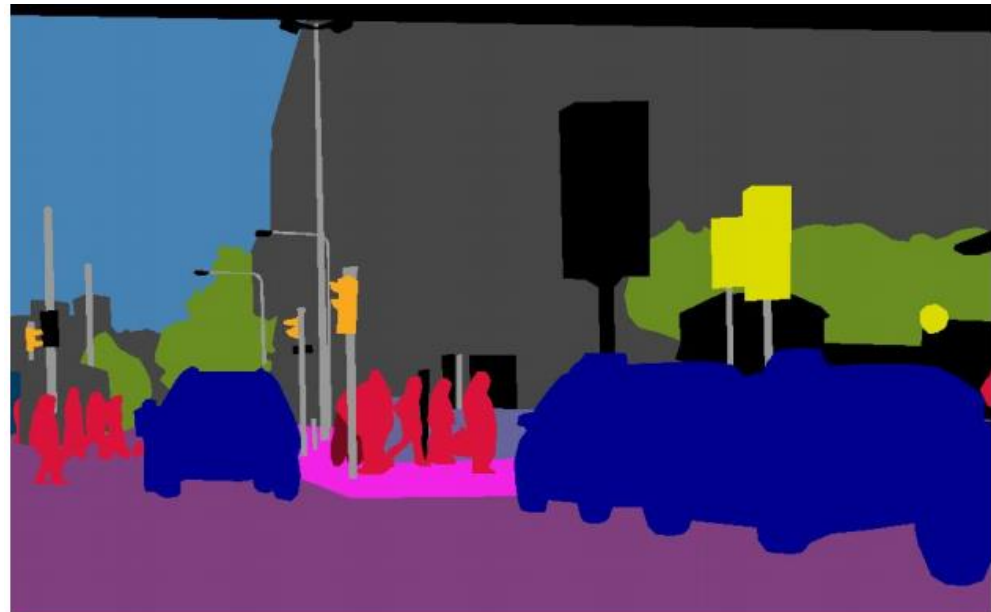
Mit Ausblick auf Object Detection





# Semantic Segmentation

Only the **type/class** of an object is important



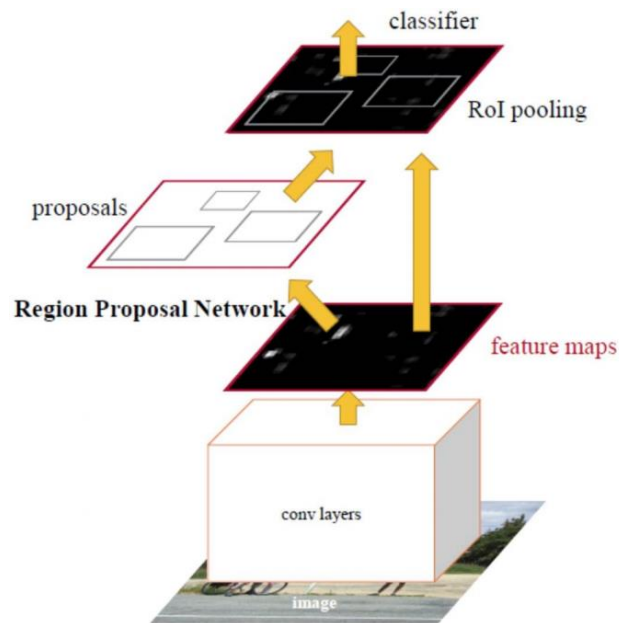
# Instance Segmentation

Differentiates between **countable** objects/classes

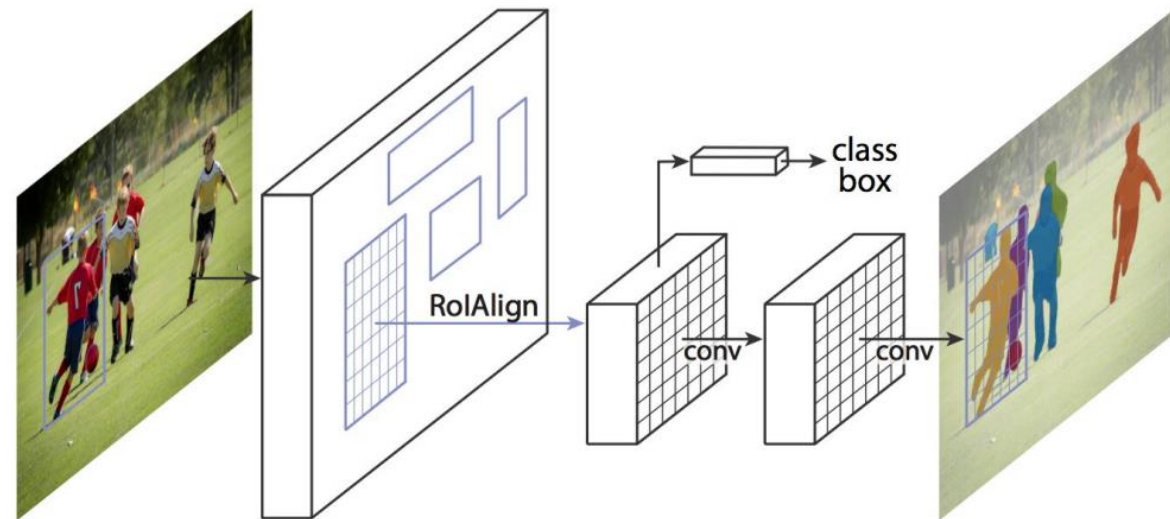


# Architecture: Mask R-CNN

- Object detection module is extended by a mask module



Faster R-CNN for object detection



Extension to Mask R-CNN



# Mask R-CNN: Example

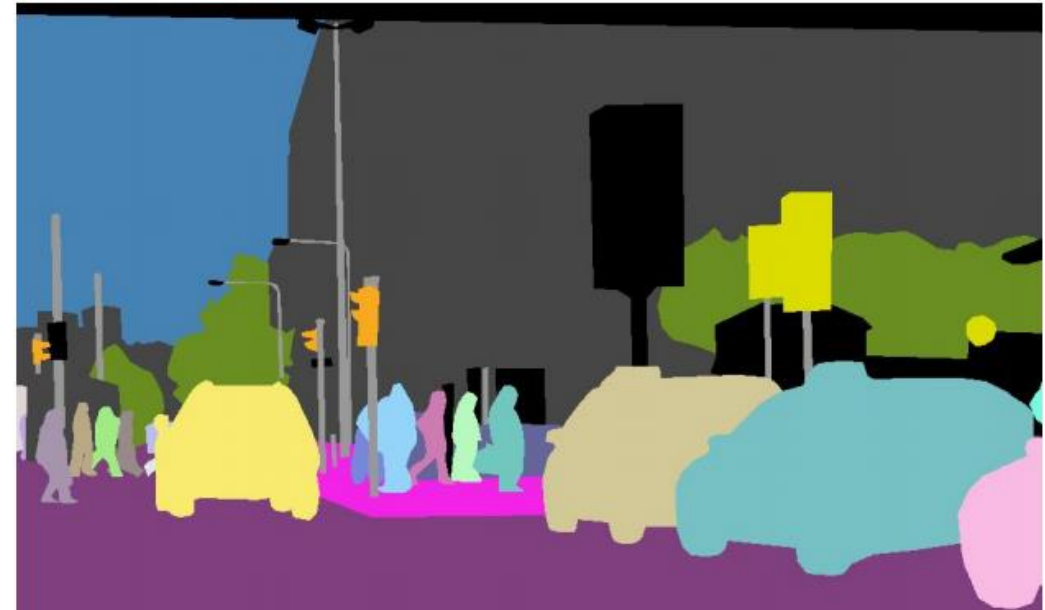
Mask R-CNN performance on COCO





# Combination: Panoptic Segmentation

Segmentation of both countable („things“) and not countable („stuff“) objects



Siehe z.B. [„UPSNET: A Unified Panoptic Segmentation Network“](#)



# Outlook

Exercise and next lecture



# Outlook: Exercise

- Investigating given FCN architectures
  - Check, if certain approaches are useful
  - Pseudo-Code
  - Compute dimensions
  - Implement (with given data)





# Outlook: Next lecture

## Object detection:

- Two big approaches
  - Two stage detectors
    - Faster R-CNN
  - One phase/single shot detectors
    - Single Shot Detector, YOLO

